

ZGCM-1: A Fully Open and Extremely Efficient Foundation Model for Math and Agentic Search

ZGCM Team^{1,2}

¹Zhongguancun Academy ²Zhongguancun Institute of Artificial Intelligence

While foundation models continue to push the frontiers of mathematical reasoning and agentic problem solving, the broader academic community has been largely excluded from this progress due to prohibitive compute requirements and closed training recipes. In this work, we present **ZGCM-1**, a fully open 7B dense foundation model trained from scratch with extreme data, system, and algorithmic efficiency. ZGCM-1 is founded on a core premise: compact models cannot passively memorize the open web, but can overcome parametric capacity limits by coupling deliberate internal thinking with active external tool use. To support this paradigm across a 256K context, we develop an end-to-end, high-efficiency open training recipe: (1) Architecture & System Co-design: interleaved gated sliding-window and full attention, and a stable FP8 Muon optimizer; (2) Progressive Curriculum & MDP Mid-Training: context scaling across 16K, 64K, and 256K, and the reformulation of interaction traces into Markov Decision Processes. Furthermore, we establish an AI-native R&D workflow where agent swarms autonomously manage cluster operations, data curation, and rapid diagnostic evaluation. Extensive evaluations show that ZGCM-1-7B is competitive across 7B model family on general benchmarks. On several challenging mathematical reasoning and agentic search suites, it remains competitive with frontier models orders of magnitude larger, such as Qwen3-235B-A22B and GLM-5.1. We also show that our pre-training design offers a $\sim 4.2\times$ efficiency improvement in 16K pre-training time-to-loss. Across the full development lifecycle, we distill eight actionable empirical findings—spanning architectural scaling, SFT quality pruning, long-context generalization, and agentic co-training dynamics. To facilitate community research, we open-source model weights from the pre-training, mid-training, and post-training stages, intermediate checkpoints, training code, per-stage data and data recipes, and W&B logs.

Code: <https://github.com/zgcagi/ZGCM-1>
Model: <https://huggingface.co/zgcagi/ZGCM-1-7B>
Data: <https://huggingface.co/datasets/zgcagi/ZGCM-1-Data>



	Average Rank	AGIEval SAT Math	AIME 2024	AIME 2025	AIME 2026	AQuA RAT	HardMath Mini	HMMT 2025	HMMT 2026	IMO-Answer Bench	MATH 500	MATH-P Hard	Olympiad Bench	ARC AGI-1	miniF2F
 ZGCM-1-7B	2.43	1	3	2	1	6	4	1	1	3	1	4	3	1	3
Qwen3-8B-Distill	2.57	4	1	1	3	5	2	2	2	1	2	2	5	5	1
MiniCPM4.1-8B	3.00	3	1	2	2	3	5	3	3	2	5	1	1	4	7
Qwen3-8B	3.14	1	4	5	4	1	1	5	4	5	3	3	2	3	3
Olmo-3-7B-Think	4.36	5	5	4	5	4	3	4	5	4	6	6	6	2	2
MiMo-7B-RL	5.43	7	6	6	6	2	7	6	6	6	3	5	4	7	5
OpenThinker3-7B	6.64	6	7	7	7	7	6	7	7	7	7	7	7	6	5

Figure 1 Per-benchmark ranks for ZGCM-1-7B and models at comparable scale across 14 reasoning benchmarks. Full results appear in Section 5.




Figure 2 Overview of ZGCM-1. Left: benchmark performance across reasoning and agentic tasks compared with 7B-scale and frontier models. Right: four key technical highlights: (1) FP8 pre-training with Muon and TWEQ achieving $\sim 4.2\times$ time-to-loss speedup, (2) hybrid sliding-window/global attention, (3) MDP mid-training that reformulates interaction traces into state-action supervision with context scaling to 256K, and (4) AI-native R&D where each researcher directs agent swarms across the full development lifecycle. BFS denotes Binary Function Search. Qwen3-8B-Distill denotes DeepSeek-R1-0528-Qwen3-8B. Full results and evaluation protocols appear in [Section 5](#).

1 Introduction

Foundation models are advancing at an extraordinary pace, pushing the frontier of long-horizon reasoning (DeepSeek-AI, 2026; Kimi Team, 2026) and tool-augmented agency in real-world environments (GLM-5-Team, 2026; Qwen Team, 2026; OpenAI, 2026; Anthropic, 2026). Despite these breakthroughs, foundational research remains encumbered by two practical bottlenecks:

- **The Scale Barrier:** Frontier reasoning and deep agentic search are widely seen as the exclusive preserve of hundred-billion-parameter systems, locking compute-constrained researchers out of training and exploring frontier-grade intelligence.
- **The Opacity Barrier:** Most competitive models are released strictly as *open-weight* rather than *fully open-source*. Upstream filtering recipes, mid-training curricula, long-context schedules, and multi-turn agent traces remain proprietary black boxes, preventing systematic study of training dynamics and capacity limits.

Our Motivation & Core Thesis. To address these barriers, we present **ZGCM-1**, a 7.39B dense foundation model trained from scratch under a transparent, open-science paradigm. We challenge the notion that advanced intelligence strictly requires massive parameter scales, centering our design on a straightforward thesis:

Compact models are inherently bounded by static parametric capacity, but they can transcend this limitation through a dual engine of deliberate internal thinking and active external seeking.

Rather than relying on passive memorization of the open web, ZGCM-1 bridges knowledge gaps by coupling long-horizon chain-of-thought reasoning with autonomous tool use—actively gathering web evidence, interacting with system terminals, and analyzing stripped binary programs.

High-Efficiency Open Recipes. To make training and inference tractable on academic compute budgets, we build an efficient full-stack pipeline across 256K contexts:

1. **Hybrid Attention Architecture:** We interleave gated sliding-window attention (SWA) with global attention at a 5:1 ratio, reducing per-token KV-cache footprint by $6.4\times$ and delivering a $3.94\times$ **throughput speedup** at 256K context over standard full attention.
2. **System-Algorithm Co-Design:** We combine the Muon optimizer (Jordan et al., 2024; Liu et al., 2025), hybrid FP8 precision, and TWEO outlier regularization (Liang et al., 2025), achieving a $\sim 4.2\times$ **pre-training time-to-loss speedup** over an AdamW/BF16 baseline.
3. **Curriculum Mid-Training with MDP Supervision:** We progressively scale context across 600B tokens (16K $\rightarrow$ 64K $\rightarrow$ 256K) while reformulating interaction traces into Markov Decision Process (MDP) state-action transitions to provide dense, step-level supervision.
4. **Execution-Grounded Alignment & Mixed SFT:** We apply mixed think/no-think fine-tuning on execution-verified trajectories balancing deep reasoning with direct-response efficiency.

Empirical Feasibility. Evaluations across 20 standard benchmarks confirm our thesis (Figure 2, Section 5):

- **Reasoning & Mathematics:** ZGCM-1-7B ranks first on average across 14 reasoning benchmarks at the 7B–8B scale while scoring **75.0% on AIME 2026**, **97.1% on MATH-500**, and **70.4% on HMMT 2025**.
- **Agentic Search & System Agency:** Deliberate thinking coupled with tool interaction allows ZGCM-1 to contend with frontier models orders of magnitude larger (e.g., Claude 4 Sonnet, Kimi-K2, and GLM-5.1), achieving **63.1% on WebWalkerQA**, **19.4% on BrowseComp**, and **62.0% on Binary Function Search**.

AI-Native R&D. We integrate researcher-directed AI agents throughout the model development lifecycle, from data processing and experimentation to evaluation and deployment. A shared agent harness combines human context artifacts with validated scripts, workflows, and debugging experience, enabling agents to execute tasks, inspect feedback, and iterate. Our Atomic Capability Evaluation suite provides rapid diagnostic feedback to guide development. Assessments from nine core contributors characterize both the benefits and limitations of this workflow: experimentation, monitoring, and deployment receive higher autonomy ratings, while architecture and learning algorithm design remain more dependent on human judgment and direction.

Empirical Findings and Open Science. Across the development lifecycle, we distill **eight empirical findings** spanning architectural efficiency, training dynamics, post-training data selection, long-context generalization, agentic co-training, and the benefits and limitations of AI-native R&D. To support community research, we release model weights from the pre-training, mid-training, and post-training stages, intermediate checkpoints, training code and configurations, data recipes, W&B logs, and evaluation harnesses at <https://github.com/zgcagi/ZGCM-1>.

2 Architecture

2.1 Model Overview

ZGCM-1 follows a decoder-only Transformer architecture with Grouped-Query Attention (GQA; Ainslie et al. 2023), RMSNorm (Zhang and Sennrich, 2019), SwiGLU activation (Shazeer, 2020), and Rotary Position Embedding (RoPE; Su et al. 2024). The model has approximately 7.39 billion parameters, 32 Transformer layers, a hidden dimension of 4,096, a SwiGLU intermediate dimension of 11,008, 32 query heads and 8 key-value heads (head dimension 128), and a maximum context length of 256K tokens.

The architecture uses a hybrid causal-attention backbone that interleaves gated (Qiu et al., 2025) sliding-window attention (SWA; Child et al. 2019; Beltagy et al. 2020) with global attention at the same 5:1 local-to-global ratio previously deployed at production scale by Gemma 3 (Gemma Team, 2025). Of the 32 layers, 27 use gated SWA with a 128-token window and five use global causal attention. The global layers are placed at layers 6, 12, 18, 24, and 30, giving five repeated blocks of five local layers followed by one global layer, plus two trailing local layers. The model further applies QK normalization (Dehghani et al., 2023), implemented with RMSNorm, and Partial RoPE with a rotary fraction of 0.33.




Figure 3 Hybrid attention architecture of ZGCM-1. The right panel shows the backbone formed by gated sliding-window and global attention layers; the left panel details the gated sliding-window attention module and its corresponding attention masks. The last global layer is layer 29 (0-indexed), followed by two trailing SWA layers.

As illustrated in Figure 3, for a normalized hidden state h , the gated SWA module forms query, key, value, and gate projections in parallel. RMS normalization and Partial RoPE are applied to the query and key branches. If $A_{\text{SWA}}(q, k, v)$ denotes the 128-token sliding-window GQA output, the gated attention output is

$$\text{GatedSWA}(h) = o_{\text{proj}}(A_{\text{SWA}}(q, k, v) \odot \sigma(g_{\text{proj}}(h))).$$

The learned sigmoid gate modulates the local-attention output element-wise before the output projection. Global-attention layers omit this gating and attend over the full context, allowing information flow beyond the local window.

2.2 Hybrid-Attention Experiments

To select the production attention schedule, we compare full attention, GQA with FlashAttention-2 (Flash-GQA; Ainslie et al. 2023; Dao 2023), MLA (DeepSeek-AI, 2024), and three SWA schedules with local-to-global ratios of 1:1, 3:1, and 5:1. All configurations are trained for 10B tokens at sequence length 4,096 on eight H100 GPUs with a common training budget.

We measure optimization quality by mean loss over the final 50 training steps and training efficiency by tokens per second per GPU. As shown in Figure 4a, SWA 5:1 achieves the highest throughput (9,566 tokens/s/GPU) while matching the tail loss of the full-attention baseline (1.93). SWA 3:1 reaches the lowest tail loss (1.92) at slightly lower throughput. MLA incurs a substantial throughput penalty (7,645 tokens/s/GPU) without a corresponding loss improvement. We adopt SWA 5:1 for the production model: it offers the highest observed throughput at a competitive tail loss. The production architecture additionally incorporates a 128-token window, QK RMS normalization, GQA, and local output gating.




Figure 4 Architecture experiments. (a) Quality-throughput trade-off among non-parameter-matched 7B attention configurations trained for 10B tokens (lower loss and higher throughput are better). (b) Throughput speedup of SWA schedules over full attention as context length grows from 4K to 256K. (c) KV cache memory footprint of three iso-parameter architectures (batch = 1, bf16): full attention (32 global layers), linear attention (GDN 3:1, 7 global + 21 linear, L=28), and our hybrid SWA 5:1 (5 global + 27 SWA).

We further compare full attention, SWA 3:1, and SWA 5:1 as context length increases from 4K to 256K on the same 7B backbone with a constant token budget per step. Figure 4b shows that the throughput advantage of SWA 5:1 over full attention widens with context length, growing from 1.13 $\times$ at 4K to 3.94 $\times$ at 256K. SWA 3:1 provides an intermediate speedup profile. This widening gap makes gated SWA well suited for long-context training and inference, where full attention becomes the dominant computational bottleneck.

2.3 KV Cache Analysis

The hybrid architecture yields substantial memory savings at inference time. Because the 27 SWA layers retain only a fixed 128-token window in their KV cache while only the 5 global layers store the full sequence, the per-token KV footprint drops from 128 KiB (full attention) to 20 KiB. Figure 4c compares three iso-parameter architectures: full attention (32 global layers, 7.39B), a linear-attention baseline (GDN 3:1 with 7 global + 21 linear layers, 7.47B), and our hybrid SWA 5:1. At 256K context, full attention requires 32.0 GiB of KV cache, GDN 3:1 requires 7.0 GiB, and our hybrid requires only 5.0 GiB (a 6.4 $\times$ reduction over full attention). Since autoregressive decoding is memory-bandwidth-bound, this smaller KV footprint directly translates to higher decode throughput, making the architecture particularly efficient for long-reasoning and agentic-search workloads that routinely operate at long context lengths.

3 Pre-Training

Pre-Training consists of two consecutive phases. As illustrated in Figure 5, General Pre-Training builds broad language, knowledge, mathematics, and code capabilities across two data stages. Mid-Training then retains the full-sequence causal language-modeling objective while introducing denser reasoning, instruction, and agentic data and progressively extending the context from 16K to 64K and 256K. Detailed corpus construction, curriculum experiments, and training configurations are reported in Appendix A.

3.1 General Pre-Training

3.1.1 Data Mixture

We first search for data mixtures with a 0.3B-parameter proxy model trained on approximately 30B tokens. Training loss and evaluation signals for knowledge, code, and mathematics guide iterative adjustments to candidate mixtures, reducing the cost of mixture exploration at the target model scale, following the broader practice of proxy-model mixture optimization (Liu et al., 2024; Team Olmo, 2025). Based on the mixture-search results, we balance capabilities across domains against loss-convergence efficiency to define the final training data recipe.



Figure 5 General Pre-Training is organized into two data stages, followed by three Mid-Training stages with progressively longer context lengths.

Web data. Curated English web corpora (Wang et al., 2025) form the main source, complemented by a controlled allocation of Chinese web data. We retain upstream document-level quality scores and quality strata produced by validation-driven filtering, and normalize accepted text and language metadata into a common document representation.

Academic and OCR data. We combine educationally filtered PDF views (Kydliček et al., 2025) with OCR-derived scientific content released by Ai2 (Allen Institute for AI, 2025), produced with the olmOCR pipeline (Poznanski et al., 2025). For scientific papers collected directly from arXiv, our in-house pipeline extracts and cleans the text while retaining the technical content required for training.

Code data. The code mixture covers both code-rich web pages and open-source repository files (NVIDIA Corporation, 2025a,b). We convert structured releases into text views, normalize file-level metadata, and clean repository artifacts and non-source content while preserving programs, technical documentation, and explanatory code text.

Mathematics data. We prioritize higher-quality tiers from mathematical corpora (Zhou et al., 2026), together with classifier-filtered web mathematics and textbook collections following the math data recipe of SmolLM2 (Allal et al., 2025). Processing combines heuristic cleaning, quality-model selection, and formula-preserving normalization so that LaTeX expressions remain embedded in their surrounding reasoning context, as in Proof-Pile-2 (Azerbayev et al., 2023).

LaTeX papers. We combine the arXiv LaTeX-source slice of RedPajama-1T (Together Computer, 2023) with filtered recent TeX sources. The accepted views recover the main textual stream, filter malformed or content-poor documents, and preserve equations and scientific document structure.

Specialized reasoning data. Reasoning-oriented material is maintained as a separate source family rather than folded into the general web or mathematics pools. We draw on Nemotron-Pretraining-Specialized-v1 (NVIDIA Corporation, 2025c), validate source schemas and versions, select the designated high-quality reasoning views, and admit them to Stage 2 as an independently controlled mixture component.

Each source family undergoes source-specific language and quality filtering, together with text extraction or repository cleaning where required. Accepted documents are normalized into a common record schema, materialized as versioned shards, and registered in source manifests. We then tokenize and index the shards with the GLM-5.1 tokenizer. Cross-stage deduplication excludes all Stage-1 content from Stage-2 selection before each source is sampled according to the final data recipe.

The final General Pre-Training corpus contains approximately 0.99T tokens in Stage 1 and 3.20T tokens in Stage 2. As shown in the upper donuts of Figure 6, web data remains the largest source in both stages. Stage 2 increases the relative contribution of code and mathematics, introduces a specialized reasoning component, and retains broad coverage of academic, OCR, and LaTeX content.

3.1.2 Curriculum Pretraining

Recent work shows that ordering training examples by difficulty rather than sampling randomly can improve pre-training efficiency within a fixed token budget (Zhang et al., 2026). We use the 0.99T-token Stage 1 as a curriculum pretraining phase. General-language documents are presented from lower to higher lexical complexity, while code and mathematics are interleaved separately. The controlled comparison and qualitative examples are reported in Appendices A.4 and A.4.1. We find that lexical-complexity ordering is cheap to compute at corpus scale for general-language data. However, it does not reliably reflect the intrinsic difficulty of code or mathematics. We therefore apply the ordering only to non-code, non-mathematics data,

remove extreme high-complexity outliers that are usually corrupted or garbled text, and interleave code and mathematics independently. In the 7B probe, this schedule lowers coding BPB from 1.99 to 0.81 and mathematics BPB from 0.97 to 0.94, while general-benchmark BPB rises by 0.03 to 0.09. The mathematics comparison uses matched evaluation sampling, while the coding gap indicates direction rather than a matched effect size. The ordering applies only to Stage 1, and later stages train on the full mixture without complexity ordering.

Finding 1 (Curriculum Pre-training): Lexical complexity is a cheap corpus-scale ordering signal for general-language text, but a poor difficulty proxy for code and mathematics, where surface statistics track boilerplate rather than algorithmic or reasoning depth. The resulting schedule, which sorts general-language documents while interleaving code and mathematics independently, improves BPB on both technical domains and trades 0.03–0.09 BPB of general-domain fit for gains concentrated where the model is targeted.

3.1.3 Hyperparameters

Optimization. We use NVIDIA Megatron Core (NVIDIA Corporation, 2026) for distributed training on H100 GPUs. Matrix parameters are optimized with Muon (Jordan et al., 2024; Liu et al., 2025) (momentum 0.9, spectral scaling, five Newton–Schulz steps, constant learning rate 2×10^{-4} , weight decay 0.1, gradient clipping 1.0). Scalar parameters use Adam.

Numerical precision. Matrix multiplications use Transformer Engine hybrid FP8 (E4M3 forward, E5M2 backward) (Micikevicius et al., 2022) with delayed scaling; scaling factors are updated from the maximum absolute value over a 1,024-step history. All other operations retain BF16 or FP32 precision. TWEO (Liang et al., 2025) is applied as an activation regularizer to suppress extreme intermediate values, complementing the delayed-scaling rule. The 16K production run sustains approximately 585 model TFLOP/s/GPU, i.e. approximately 60% BF16-equivalent MFU against the 989 TFLOP/s H100 BF16 dense peak. Full parallelism configurations are reported in Appendix A.1.

Finding 2 (TWEO enables FP8 Pretraining): Suppressing intermediate activation outliers is crucial for FP8 training stability; combining delayed scaling with TWEO (Liang et al., 2025) regularization prevents numerical divergence while sustaining 60% BF16-equivalent MFU at scale.

Training efficiency. We estimate 16K pre-training time-to-loss relative to a comparable OLMo 3-style 7B BF16/AdamW baseline (Team Olmo, 2025). Four factors contribute: SWA 5:1 yields a $1.4\times$ throughput gain over full attention at 16K (Section 2.2, Figure 4); the FP8 precision-and-systems configuration contributes approximately $1.5\times$; Muon provides approximately $1.8\times$ step-to-loss efficiency over AdamW; and the Pre-LN contributes an estimated $1.1\times$ data efficiency from exploratory 7B evidence. Multiplying these gives

$$1.4 \times 1.5 \times 1.8 \times 1.1 \approx 4.2.$$

This roughly indicates a fourfold improvement in 16K pre-training time-to-loss.

Finding 3 (Pretraining Efficiency Gain): The co-design of hybrid SWA ($1.4\times$), FP8 mixed-precision systems ($1.5\times$), the Muon optimizer ($1.8\times$), and Pre-LN ($1.1\times$) yields a cumulative $\sim 4.2\times$ speedup in 16K pre-training time-to-loss over a standard BF16/AdamW baseline.

3.2 Mid-Training

3.2.1 Data Mixture

Mid-Training is capability-oriented continued pre-training. We apply pool-wide deduplication to improve token efficiency and control repeated content (Lee et al., 2022). From the resulting 2.86T-token candidate pool, we construct a 600.51B-token sampled schedule organized into 16K, 64K, and 256K context stages.




Figure 6 Data mixture composition shown as donut charts for the two General Pre-Training stages and the three Mid-Training context stages. Wedge areas represent each category’s share of the corresponding stage’s total tokens.

The maximum sequence length increases progressively across these stages, following the multi-stage context expansion of Qwen2.5-1M (Yang et al., 2025b); other open technical reports extend the context in a single final pre-training stage (LLM-Core-Team Xiaomi, 2025; Yang et al., 2025a). Each later stage remains cumulative rather than replacing shorter sequences: the 64K stage contains 180.89B tokens at up to 16K and 59.11B tokens in the 16K–64K range, while the 256K stage combines 127.81B tokens at up to 16K, 21.72B tokens in the 16K–64K range, and 30.98B tokens above 64K.

The 16K stage trains at the standard context length with a balanced mixture of code, mathematics, knowledge, and reasoning. The 64K stage introduces longer documents and reasoning sequences while retaining a substantial share of shorter-context data. The 256K stage adds ultra-long documents, cross-document information, and long-horizon agentic trajectories, and continues to replay data from the 16K and 64K length ranges to preserve coverage of conventional-context tasks.

The lower donuts of Figure 6 summarize the top-level capability mixture. Code and mathematics each remain close to 20% throughout the curriculum, preserving core programming and reasoning capabilities. As the context length grows, knowledge data increases from 10.50% to 14.23%, and agentic data increases from 1.50% to 3.30%, raising the density of long-document, tool-interaction, and multi-step task examples. Pre-training replay decreases from 13.00% to 9.00% but remains present to maintain broad coverage as higher-density data is introduced. Web, QA, reasoning, and instruction data remain comparatively stable across the three stages. Exact length accounting and construction details are given in Appendix A.5.




Figure 7 Mid-Training data processing pipeline.

3.2.2 Reasoning Data

Reasoning sources are normalized to a common schema and routed into accepted, rewrite, pending, holdout, or rejected views according to self-containment, answer evidence, and reasoning value. Rule-cleaned examples form the main pool.

Valuable questions whose original traces are unsuitable for training are extracted as standalone problems, reconstructed or improved by teacher models, and admitted only after checks for parsing validity, completeness, template and source leakage risks, duplication, and trainability. This separates question value from the quality of an upstream reasoning trace. Details are reported in [Appendix A.5.1](#).

3.2.3 Instruction Data

We collect a large-scale instruction corpus that includes Tulu- and FLAN-derived data ([Lambert et al., 2024](#); [Wei et al., 2022](#)). These examples are converted into pre-training-compatible raw context. The pipeline preserves legitimate short-label tasks, routes translation-risk subsets through model review, and rewrites selected single-turn examples as multi-turn discussions. These transformations increase coverage of conversational state and sequential reasoning while retaining the full-sequence training objective. Further construction details are provided in [Appendix A.5.1](#).

3.2.4 Agentic Data

Agentic data for Mid-Training covers both general interaction traces and trajectories collected from software-engineering tasks. We retain high-quality complete interactions and, following agentic continual pre-training ([Su et al., 2025](#)), reformulate general traces as Markov decision process-style state-conditioned next-action prediction examples, combining full-trajectory context with denser supervision for local decision making. For software-engineering trajectories, we apply execution-aware filtering and preserve the interaction context needed to learn planning, tool use, and iterative refinement. Detailed construction procedures, filtering criteria, and corpus statistics are provided in [Appendix A.5.2](#). For each source family, we use the agent-driven self-iterating governance process shown in [Figure 8](#) to develop and validate a source-specific processing recipe.

AI - DRIVEN SELF- ITERATING DATA GOVERNANCE

From source routing to a validated reasoning-data pool




Figure 8 AI-driven self-iterating data governance pipeline. Each source family is routed and sampled independently, then processed through source-specific rule filtering, strong-model audit, failure mining, human spot checks, iterative script revision, and held-out validation before accepted examples enter the validated reasoning-data pool.

3.2.5 Hyperparameters

All Mid-Training stages use full-sequence causal language modeling over raw context (distinct from the assistant-only supervision used in SFT, as detailed in [Section 4.1.3](#)). We maintain approximately 12.6M tokens per optimizer step across all three stages by adjusting global batch sizes (768, 192, and 48 for 16K, 64K, and 256K respectively). The learning rate is 10^{-4} , and the 256K stage uses a RoPE base of 10M with full activation recomputation. We evaluate both a direct 256K route (30B tokens) and a staged route (10B at 64K then 20B at 256K); the staged route achieves a slightly lower final loss. Full configurations are in [Appendix A.1](#).

The resulting C1 staged run is summarized by its loss trajectory and learning-rate schedule in [Figure 9](#).

3.3 Monitor During Training

Throughout General Pre-Training and Mid-Training, we continuously monitor training dynamics, numerical stability, hardware and systems status, training efficiency, and resource utilization.

To monitor capability development during General Pre-Training, we directly evaluate saved base checkpoints with a fixed capability suite covering knowledge, question answering, mathematics, reasoning, coding, and basic symbolic operations. No supervised fine-tuning or other post-training adaptation is applied before evaluation. In parallel, we monitor perplexity, bits per byte (BPB), normalized loss, key-token logits, and optimizer-related statistics that reflect training dynamics. These signals help diagnose optimization and systems issues, detect capability regressions, and inform adjustments to data mixtures, curricula, and stage transitions.

Checkpoint-level evaluations reveal heterogeneous capability development across domains. Knowledge and question-answering benchmarks generally improve over the monitored training window, but the rate and timing of improvement differ across benchmarks. Mathematics and coding capabilities also improve, although their gains are concentrated in different intervals and accompanied by larger short-term fluctuations. We therefore treat these curves as descriptive diagnostics of domain-dependent learning dynamics rather than as evidence of a universal monotonic trend. Detailed trajectories for all benchmarks are provided in [Appendix A.2](#).

As an example, continuous throughput monitoring after resuming a training job revealed a gradual decline across matched windows of approximately 80 steps, from 585 to 554 model TFLOP/s/GPU. Releasing overlap buffers reduced but did not eliminate the decline. Adding periodic CUDA allocator cache clearing and garbage collection every 100 iterations restored stable throughput at 585 model TFLOP/s/GPU. This fix was deployed into the production training loop.




Figure 9 Token-level cross-entropy (CE) loss (blue, left axis) and learning-rate schedule (terracotta, right axis) for the staged Mid-Training run. The solid blue curve is a 200-point moving average, while the light-blue curve shows raw loss values. Both metrics share the cumulative-token axis; vertical dotted lines mark transitions between the 16K, 64K, and 256K context stages.

A complementary short-SFT probe evaluates successive training milestones after lightweight supervised adaptation. Unlike the direct base-checkpoint evaluation above, this probe measures the downstream capabilities that can be elicited after supervised adaptation. Its results are reported in [Appendix A.3](#).

4 Post-Training

The post-training pipeline consists of supervised fine-tuning (SFT) followed by mixed reinforcement learning (RL). SFT establishes response modes with general and agentic data, while RL yields capability gains on single-domain tasks such as mathematics and code. Additional data and implementation details are provided in [Appendix B](#).

4.1 Supervised Fine-Tuning

[Figure 10](#) summarizes our SFT data curation pipeline. We organize the corpus into general and agentic data, with general data comprising general instruction and reasoning examples. Before mixture selection, all examples are normalized to a common message and tool schema and filtered for structural validity, response correctness, and tool-turn consistency. Candidate-run evaluations then determine a mixture that balances capability coverage, long-reasoning supervision, and joint general-agentic training. Finally, we remove benchmark contamination and convert the retained examples into packed sequences with assistant-only loss masks and length-based routing.

4.1.1 General Data

Our SFT corpus combines open-source datasets with internally distilled data and contains 4,921,933 examples organized into general and agentic data. As shown in [Table 1](#), general data covers instruction following, knowledge, mathematics, science, code, dialogue, and reasoning, while agentic data targets interaction with tools and external environments.

Data Quality Filtering and Tiered Curation. We establish a multi-stage data curation pipeline combining deterministic heuristic rules with model-based quality scoring. In the initial stage, rule-based checks systematically eliminate structural defects, including missing dialogue turns, invalid role assignments, duplicate generations, prompt template leakages, and malformed reasoning traces. Following rule-based filtering, we deploy a model-based evaluator calibrated against human-annotated pilot benchmarks to grade candidate instances along multiple pedagogical and technical axes: *educational value*, *logical soundness*, *step-by-step*




Figure 10 Supervised fine-tuning data curation pipeline. General data comprises general instruction and reasoning examples, while agentic data forms a separate branch. All sources undergo schema formatting, filtering and verification, mixture selection, benchmark decontamination, and preparation as packed training sequences with assistant-only loss.

Table 1 Composition of our supervised fine-tuning corpus.

Component	Row share	Primary role
General data	96.46%	Instruction following, knowledge, mathematics, science, code, dialogue, and reasoning
Agentic data	3.54%	Deep research, software engineering, and terminal interaction through structured tools
Total	100.00%	All supervised fine-tuning data

reasoning coherence, factual consistency, and safety compliance. Rather than applying rigid binary thresholds, fixed multi-criteria scoring rules classify examples into discrete quality tiers. High-scoring instances offering dense supervision and diverse capability coverage are prioritized, borderline cases are routed to auxiliary verification pipelines, and noisy, uninformative, or high-risk entries are strictly discarded. Refusal samples undergo dedicated isolation and filtering to eliminate over-refusals and environment-dependent artifacts while preserving essential safety alignment. Empirical comparisons across candidate training runs indicate that aggressive tiered filtering—pruning approximately 50% of the raw SFT candidates—outperforms training on the uncurated full corpus in aggregate, though individual benchmarks do not all move in the same direction. A controlled three-way comparison of filtering strength is reported in [Appendix B.1.1](#).

Finding 4 (Data Quality Trumps Volume in Posttraining): SFT quality matters more than data volume. In controlled comparisons, pruning roughly half of the candidate pool improved the six-benchmark mean from 67.78 to 68.83, with gains concentrated in reasoning (BBH ([Suzgun et al., 2023](#)) +10.10) and losses confined to a subset of tasks. We therefore prioritize reliable, high-value supervision over maximizing the number of retained examples.

Decontamination. To prevent benchmark contamination and ensure reliable generalization metrics, we enforce a strict cross-benchmark decontamination protocol. The entire candidate SFT corpus is scrubbed against all downstream evaluation and validation sets using an 8-gram matching sliding window. Any training sample exhibiting greater than 50% 8-gram overlap with any evaluation prompt or reference response is permanently purged from the corpus.

Capability Balancing and Reasoning Mixture. To determine the optimal data composition across domain capabilities, we conduct controlled mixture ablation sweeps on candidate SFT models across mathematics, code generation, multi-turn dialogue, and complex instruction following. These sweeps reveal that disproportionately scaling long chain-of-thought (CoT) trajectories introduces verbosity bias and impairs instruction adherence.

We systematically perform iterative grid calibrations to harmonize the proportions of long-form reasoning, direct-response QA, and strict formatting directives, arriving at a balanced Pareto-optimal mixture.

Finding 5 (Long-CoT Trade-offs and Calibration): Long chain-of-thought (CoT) supervision exhibits non-monotonic utility: oversaturating the training mixture with extended reasoning trajectories degrades general instruction-following performance. Dynamically calibrating data mixture ratios mitigates this degradation, preserving strong instruction adherence without compromising long-CoT reasoning gains.

Unified Think and Direct-Response Supervision. The resulting SFT corpus deliberately interweaves *think* examples (containing explicit, end-to-end reasoning chains encased in structured reasoning tags) with *no-think* examples (featuring concise, immediate responses). This dual-mode design equips a single set of model weights to dynamically toggle between full deliberation and efficient zero-shot responses depending on user-specified system instructions or runtime budgets. Crucially, intermediate checkpoint evaluations confirm that joint training induces positive cross-modal transfer: exposure to structured reasoning trajectories directly sharpens the accuracy of no-think responses across coding, mathematics, and logical reasoning benchmarks (detailed in [Appendix B.4](#)).

4.1.2 Agentic Data

Our agentic SFT data comprises three parallel branches: deep research, software engineering, and terminal interaction. During data construction, we align each task family’s training schema with its downstream inference environment, including system instructions, message roles, tool definitions, structured calls, observation placement, and final-answer conventions. Our experiments show that schema mismatches substantially degrade downstream agent performance. The Deep Research schema is documented in [Appendix B.2.1](#). Within these aligned schemas, multi-turn action-observation trajectories supervise tool selection, information gathering, environment interaction, and final response generation.

Tool-Protocol Validation. Across agentic branches, we convert source-specific and legacy tool-call formats into the target structured representation. We validate argument structure, tool-call-observation pairing, unique call identifiers, and complete assistant termination, and remove textual instructions that conflict with the target protocol. Automatic repairs are limited to examples with unambiguous call-result correspondence; irrecoverable examples are discarded. Stable example identifiers preserve the links among source records, quality decisions, and protocol revisions.

Agentic behavior also depends on instruction following, reasoning, and domain knowledge. We therefore compare a sequential schedule that first trains on general-capability data and then continues on agentic data with a joint schedule that interleaves general and agentic examples throughout SFT. The joint schedule yields stronger downstream performance and is used for the final training run.

Finding 6 (The Necessity of General Data for Agent Training): Agent-only fine-tuning leads to degraded interaction fidelity, as agentic proficiency inherently relies on broad generalist capabilities. Co-training with general instruction data provides indispensable grounding in step-by-step reasoning and precise constraint following, yielding substantially superior agentic performance.

Deep Research Trajectories. We collect 20,217 multi-step research trajectories and normalize them into a shared interaction environment with `search` and `visit` tools. Structured tool calls and environment observations follow a common representation. We remove trajectories with invalid arguments, unmatched action-observation pairs, or malformed tool returns. Each accepted trajectory is provided in aligned think and no-think forms.

Software-Engineering Trajectories. We construct the software-engineering branch from collected repository-interaction trajectories. The execution-grounded subset covers repository interaction in executable environments, whereas the execution-free subset provides supervision for repository understanding, file localization,

tool selection, and patch planning. After subset-specific filtering and scoring, we retain 30,014 execution-grounded and 30,000 execution-free trajectories and project them to the GLM-5.1 structured-tool format. The construction pipeline is detailed in [Appendix B.2.2](#).

Terminal Trajectories. We collect terminal-agent trajectories from terminal task environments through multiple independent collection routes. Each task runs in an isolated Docker environment with a persistent shell and a structured `bash` tool, followed by an environment-side verifier. We deduplicate trajectories within each collection route at the task-environment level and normalize source-specific reasoning and tool protocols. For a later controlled Agent-SFT experiment, we materialize a paired-source view containing 22,309 trajectories and a 15,748-trajectory verifier-successful subset. This view is an ablation artifact and is not the terminal subset used in the main SFT run. Further details are provided in [Appendix B.2.3](#).

4.1.3 Training Objective and Hyperparameters

SFT renders all examples in a common GLM-5.1 conversation format and applies assistant-only supervision. Depending on the mode, the target contains either an explicit reasoning trajectory followed by a final answer or a direct answer without an exposed reasoning trace. Assistant reasoning, responses, and structured tool calls contribute to the loss; system and user messages and tool observations remain masked context. This mixed think/no-think objective trains reasoning, response generation, and action selection without teaching the model to reproduce environment feedback. Packing and loss-mask validation are detailed in [Appendix B.3](#).

We train two length-specific SFT variants with maximum sequence lengths of 65,536 and 262,144 tokens. The 64K variant is trained for eight epochs, while the 256K variant is trained for ten epochs over 19.46B packed tokens and is used for the released model. Both runs use a global batch size of 48. We apply layer-wise Muon ([Liu et al., 2025](#)) to matrix parameters and Adam to scalar parameters, with the learning rate cosine-decayed from 1×10^{-4} to 1×10^{-6} , no warmup, a weight decay of 0.01, and gradient clipping at 1.0. The corresponding optimization dynamics are shown in [Figure 11](#).




Figure 11 Training dynamics of the two length-specific SFT variants. The 64K run (left) and the 256K run (right) exhibit similar loss-reduction patterns: a rapid initial decrease followed by stepwise convergence as the learning rate follows a cosine decay schedule. Curves are shown from optimizer step 50 to focus on steady-state training dynamics. Solid blue curves show a 160-step moving average of token-level cross-entropy (CE) loss, light-blue curves show raw loss values, and salmon curves show the learning rate. Both panels use the same CE-loss scale. The released model uses the 256K SFT variant.

Finding 7 (Long Agentic Mid-Training Bypasses Ultra-Long SFT): The long-context foundation established during Mid-Training allows robust reasoning and agentic behaviors to be reliably elicited at the 256K scale using only moderate-length SFT data. This effectively bypasses the need for expensive 256K-long reasoning trajectories during post-training, demonstrating that shorter-sequence supervision suffices to activate and extrapolate long-sequence capabilities established in the mid-training.

4.2 Mixed Reinforcement Learning

4.2.1 RL Data Curation

We construct a mixed reinforcement-learning corpus covering mathematics, code, and general-capability tasks. We estimate prompt difficulty through rollouts from the initial policy and remove problems that the model already solves frequently, concentrating training on prompts that continue to provide useful learning signals.

4.2.2 Reward Design

Rewards follow the evaluation semantics of each domain. Mathematics uses binary answer-correctness rewards, while code receives rewards based on the fraction of executable tests passed. General-capability tasks use corresponding correctness or instruction-following criteria. Invalid or truncated responses do not receive positive rewards. Under a maximum response length of 64K tokens, we further explore combining outcome-based rewards with the reference-policy KL regularization of GRPO (Shao et al., 2024) and a mild length penalty to stabilize long-horizon optimization. The KL term controls policy drift, while the length penalty reduces reward noise from overlong or truncated responses (Yu et al., 2025), supporting reliable multi-step reasoning without encouraging unnecessary response expansion.

4.2.3 Hyperparameters

Optimization. We optimize the policy with Group Relative Policy Optimization (GRPO) (Shao et al., 2024). For each problem, we sample multiple responses and construct relative advantages within the corresponding response group. Groups with zero reward variance do not provide effective relative learning signals; dynamic sampling therefore filters and replaces them during training (Yu et al., 2025). The actor learning rate is 2×10^{-6} .

Sampling and rollouts. Across experiments, we explore two sampling scales. The first samples 24 problems per training step and generates 16 responses per problem, yielding $24 \times 16 = 384$ trajectories. The second samples 384 problems and generates 8 responses per problem, yielding up to $384 \times 8 = 3,072$ trajectories per step. To support long-horizon reasoning, rollouts allow up to 65,536 newly generated tokens within a total prompt-response context of 98,304 tokens; the actor microbatch uses the same token budget.

5 Evaluation

5.1 Setup

We evaluate ZGCM-1 in think mode. Reported results use the released 256K SFT checkpoint (Section 4.1). Non-agentic and web-environment deep-research runs use a sampling temperature of 1.0 and top- p of 1.0. Their total context budget is 262,144 tokens, comprising up to 4,096 prompt tokens and up to 258,048 generated tokens. For non-agentic benchmarks, we report mean pass@1 over 32 runs unless the benchmark specifies another aggregation rule. For DeepSeek-R1-0528-Qwen3-8B and MiniCPM4.1-8B, GPQA-Diamond, HumanEval+, MBPP+, and LiveCodeBench v6 are four-run means; MMLU is a single run and IFEval uses Prompt Strict accuracy.

The non-agentic suite covers 20 benchmarks across four categories. For mathematical and abstract reasoning, we evaluate MATH-500 (Lightman et al., 2023), AIME 2024/2025/2026 (Mathematical Association of America, 2026), HMMT 2025/2026 (Harvard-MIT Mathematics Tournament, 2026), AGIEval SAT Math (Zhong et al., 2023), AQuA-RAT (Ling et al., 2017), HARDMath-mini (Fan et al., 2024), IMO-AnswerBench (Luong et al., 2025), MATH-P-Hard (Huang et al., 2025), and the text-only mathematics subset of OlympiadBench (He et al., 2024), together with formal theorem proving on miniF2F (Zheng et al., 2022) and abstract grid reasoning on ARC-AGI-1 (Chollet, 2019). For code generation, we use HumanEval+ and MBPP+ (Liu et al., 2023), together with LiveCodeBench v6 (Jain et al., 2024). For knowledge, we use MMLU (Hendrycks et al., 2021) and GPQA-Diamond (Rein et al., 2024); instruction following is measured with IFEval (Zhou et al., 2023).

Web-environment deep research uses a thinking-enabled ReAct policy (Yao et al., 2023) with at most 64 search-and-read steps. Binary Function Search uses a separate Ghidra-based (National Security Agency, 2019)

Table 2 Thinking-mode benchmark results (%) for ZGCM-1-7B and comparable-scale models. Dark blue cells in bold indicate the best result; light blue cells indicate the second-best result.

Benchmark	ZGCM-1 7B	DeepSeek-R1 0528-Qwen3 8B	MiniCPM 4.1-8B	Qwen3 8B	Olmo-3-7B Think	MiMo-7B RL	Open Thinker 3-7B
Reasoning & General							
MATH-500	97.13	96.32	95.60	96.20	95.10	96.20	88.40
AIME 2024	80.62	83.33	83.33	80.00	71.60	66.67	43.33
AIME 2025	73.33	75.21	73.33	63.33	64.60	53.33	43.33
AIME 2026	75.00	69.17	71.67	66.67	66.16	56.67	40.00
HMMT 2025	70.42	61.50	52.50	43.33	43.89	40.00	23.33
HMMT 2026	59.48	51.52	46.21	45.45	43.94	39.39	21.21
AGIEval SAT Math	99.09	91.14	98.98	99.09	90.45	58.64	68.64
AQuA-RAT	90.57	90.88	91.39	92.62	90.98	91.80	81.15
HARDMath-mini	56.34	61.92	51.89	63.31	58.90	37.41	37.45
IMO-AnswerBench	53.00	55.00	54.25	45.00	45.25	41.50	26.00
MATH-P-Hard	79.21	82.53	84.41	82.08	77.42	78.14	59.50
OlympiadBench	76.06	74.75	82.75	82.47	69.55	74.89	56.37
ARC-AGI-1	3.50	1.69	2.58	3.00	3.25	0.50	1.50
miniF2F	2.87	5.12	0.00	2.87	3.28	1.23	1.23
Code							
HumanEval+	90.24	88.87	89.63	80.20	89.90	88.95	87.40
MBPP+	63.23	65.54	63.96	69.10	64.70	63.96	61.40
LiveCodeBench v6	46.86	53.57	52.14	52.20	49.26	47.42	42.43
Knowledge							
MMLU	73.88	82.09	83.51	85.40	77.80	78.39	77.40
GPQA-Diamond	47.87	60.10	47.98	59.09	49.94	54.40	53.70
Instruction Following							
IFEval	75.42	72.37	73.24	87.40	88.20	61.00	51.70

interaction protocol and decoding configuration. Both harnesses are specified in [Appendix C](#). Web baseline values other than WebDancer are taken from public technical reports and model releases; WebDancer and Binary Function Search baselines are evaluated under our shared harness.

5.2 7B Scale Model Comparison

We compare ZGCM-1 with six reasoning models at the 7B–8B scale. As shown in [Table 2](#), the comparison spans 20 benchmarks across reasoning and general capabilities, code, knowledge, and instruction following.

5.3 Agentic Evaluation

We evaluate ZGCM-1 in two agentic research settings: open-web information seeking and binary-program analysis.

Web-Environment Deep Research. Web-environment deep research requires agents to iteratively search, inspect, and synthesize external evidence. [Table 3](#) compares specialized research agents, open-weight models with tools, and proprietary systems using reported results on WebWalkerQA ([Wu et al., 2025b](#)), BrowseComp ([Wei et al., 2025](#)), and GAIA ([Mialon et al., 2024](#)), with WebDancer ([Wu et al., 2025a](#)) as the specialized research-agent baseline. ZGCM-1 obtains 63.09% on WebWalkerQA, 19.43% on BrowseComp, and 42.52% on GAIA text-only.

Binary Function Search. This benchmark evaluates whether an agent can recover a target function from a stripped ELF binary using only a behavior description. We construct the benchmark from a broad collection effort spanning thousands of open-source C/C++ projects, selecting several semantically meaningful functions from each project when suitable candidates are available. The current validated data pool contains more than 11,000 tasks from hundreds of projects. We use a 50-task subset for evaluation, formed by randomly

Table 3 Web-environment deep-research results (%). Dark blue cells in bold indicate the best result; light blue cells indicate the second-best result. Dashes denote unavailable reported results.

Model / System	WebWalkerQA	BrowseComp	GAIA text-only
<i>Our Model</i>			
ZGCM-1-7B	63.09	19.43	42.52
<i>Open-source Research Agent</i>			
WebDancer-32B	38.40	2.50	40.70
<i>Open-weight Models with Tools</i>			
Qwen3-235B-A22B	59.60	2.30	45.60
Kimi-K2	63.00	14.10	57.30
DeepSeek-R1	10.00	2.00	31.10
QwQ-32B	4.30	0.50	22.30
<i>Proprietary Models with Tools</i>			
GPT-4o	33.80	1.90	34.60
GPT-5	—	54.90	76.40
Claude 4 Sonnet	61.70	12.20	68.30

sampling five tasks from each of 10 representative projects (tmux, tree, GNU Wget2, XZ Utils, YAJL, zlib, libuv, libgit2, nm, and Lua). All 10 projects are held out from the training data and do not appear in any training split; consequently, no evaluation task is derived from a project seen during training. We plan to publicly release the Binary Function Search dataset to support reproducible research on tool-assisted binary analysis.




Figure 12 Binary Function Search workflow and Ghidra command interface. The agent iterates between candidate exploration and evidence refinement, then submits an exact function-entry address for hidden-oracle validation.

As shown in Figure 12, the agent receives a behavior description containing clues about control flow, callees, constants, error handling, or outputs. Through a structured Ghidra-backed interface, it enumerates functions, searches literals and error messages, decompiles candidates, follows references, and inspects assembly. The agent iteratively narrows the candidate set by comparing this evidence with the description, then submits the exact function-entry ELF virtual address. We validate the submitted address against a hidden oracle constructed from the corresponding unstripped binary. A submission is judged correct only if it precisely matches the exact function-entry ELF virtual address; incorrect or missing submissions are judged incorrect. Further construction and scoring details appear in Appendix C.2. ZGCM-1-7B achieves 62% accuracy, far exceeding similarly sized baselines (12% for Qwen3-8B and 0% for the others). It remains competitive with larger frontier models, approaching GLM-5.1 (66%) and outperforming DeepSeek-R1 and GPT-4o.

Table 4 Binary Function Search results on 50 tasks. Valid submissions counts answers in the required format; Correct counts exact oracle matches; Accuracy is Correct/50. The dark blue cell in bold and the light blue cell indicate the best and second-best accuracy.

Model / System	Valid submissions	Correct	Accuracy (%)
<i>Larger frontier models</i>			
DeepSeek-V4 Flash	50/50	38/50	76.00
Qwen3.5-397B-A17B	50/50	38/50	76.00
GLM-5.1	50/50	33/50	66.00
ZGCM-1-7B	50/50	31/50	62.00
Kimi-K2	50/50	31/50	62.00
DeepSeek-R1	46/50	18/50	36.00
GPT-4o	50/50	9/50	18.00
<i>Comparable-scale models</i>			
ZGCM-1-7B	50/50	31/50	62.00
Qwen3-8B	50/50	6/50	12.00

6 AI-Native Research and Development

6.1 AI Assistance Across the R&D Lifecycle

We adopt an AI-native methodology throughout the entire model development lifecycle. As illustrated in Figure 13, each researcher directs LLM-based agents that contribute across all stages of development, from data engineering and architecture design to learning algorithms, experimentation, infrastructure, evaluation, and deployment. Researcher discussions produce human context artifacts (conversations, meeting notes, decisions, plans), while completed agent work produces reusable experience (validated scripts, workflows, checklists, debugging history). Both streams feed into a shared agent harness that equips every agent with the context, skills, and tools needed for effective autonomous operation.

Data cleaning. AI agents participate in the full data-processing pipeline. Given a set of quality requirements, an agent first writes heuristic filtering and transformation scripts, then executes them on the target corpus. After processing, the agent inspects the resulting data composition and distribution statistics to verify whether the output meets the predefined criteria. If discrepancies are found, the agent iterates on the processing scripts automatically, refining rules and thresholds until the data distribution satisfies all requirements.

Cluster management. Data processing, model training, and inference all run on a unified cluster platform where users submit jobs, inspect logs, and manage development machines through a single web interface. To enable AI agents to operate this platform autonomously, we iteratively co-developed a comprehensive skill package with our agents. The package exposes structured interfaces for creating development machines, submitting training jobs, modifying resource configurations, streaming logs, and filtering status fields. With these skills, agents can independently launch experiments, monitor their progress, retrieve and diagnose failures from logs, and adjust hyperparameters or resource allocations without human intervention.

Auto research. Current LLM-based agents possess sufficient long-horizon task capability to conduct semi-autonomous research when equipped with literature search and cluster management tools. Agents can explore algorithmic alternatives and hyperparameter configurations, launch training runs, analyze results, and iterate without continuous human supervision. We routinely use this auto-research workflow during ZGCM-1 development. For example, agents profiled the data-transfer characteristics of our storage cluster and made targeted adjustments to the training framework’s I/O pipeline, yielding substantial training-speed improvements entirely through autonomous agent work.

Evaluation. Model evaluation is part of our AI-native R&D workflow, serving as a development tool rather than only a final-stage assessment. We use Atomic Capability Evaluation (ACE), a fine-grained diagnostic benchmark that rapidly characterizes model capabilities and feeds back into training iterations. ACE organizes its 2,503 probes into 183 atomic capabilities across 18 categories. Under our standard setup, a complete ACE run takes approximately two to three minutes, giving a first-order profile of strengths and weaknesses at the




Figure 13 AI-native R&D workflow. Each researcher directs agents across seven core stages of model development: data engineering, model architecture design, learning algorithm design, experimentation & monitoring, infrastructure engineering, evaluation, and deployment engineering. Completed agent work produces reusable experience (validated scripts, workflows, checklists, debugging history), and researcher discussions produce human context (conversations, meeting notes, decisions, plans); both feed into a shared agent harness that equips all agents with accumulated context, skills, tools, memory, orchestration, and verification capabilities.

category, capability, and probe levels. This makes ACE well suited to frequent evaluation during ablation studies: models trained with different data mixtures, objectives, hyper parameters, or inference configurations can be compared efficiently, and regressions can be localized to affected capabilities. ACE thus closes the loop between evaluation and development, identifying which interventions help, which introduce regressions, and where further adjustments are needed. We construct the benchmark through iterative human-LLM collaboration. LLMs assist with capability decomposition, probe generation, boundary-case discovery, and scorer implementation, while researchers define capability boundaries, audit generated artifacts, and approve their inclusion. Benchmark execution is fully automated: deterministic programs score rule-expressible probes, while a separate model applies predefined rubrics when correctness cannot be expressed reliably as an executable rule. [Appendix D](#) details the taxonomy, scoring policies, validation, and results.

Human context. Researchers and developers engage in frequent project discussions that, although information-sparse, contain critical design decisions, action items, and technical rationale. We integrate our self-developed agent platform, ZGent, into the team’s regular meeting channels. ZGent records and distills the content of each discussion, accumulating context over time and automatically dispatching resulting tasks to the appropriate agents for execution. This closes the gap between human deliberation and agent action, significantly reducing the overhead of translating meeting conclusions into concrete engineering work.

6.2 AI4AI Autonomy Across the R&D Lifecycle

AI4AI (AI for AI) refers to using AI systems to build, evaluate, and improve other AI systems. To assess how much responsibility AI agents can assume across the R&D lifecycle, nine core contributors each rated 11 task categories against the five-level rubric in [Table 5](#), yielding 99 ratings. The L4/L5 boundary is autonomous objective formation: L4 executes independently within human-defined objectives, whereas L5 also identifies research objectives and coordinates work across stages. The rubric is ordinal, so small differences between means are not meaningful ([Appendix E](#) gives the task-specific criteria and all individual ratings).

Table 5 Levels of AI autonomy in R&D workflows. The levels describe autonomy rather than productivity gains or output quality. L5 provides a reference for full autonomy, rather than an assertion of demonstrated capability.

Level	Name	Defining criterion
L1	Basic Assistance	Humans direct and execute the workflow; AI assists with individual steps.
L2	Partial Automation	AI executes predefined workflows; humans specify procedures and handle exceptions.
L3	Conditional Autonomy	AI adapts and manages routine workflows; consequential decisions require human approval.
L4	High Autonomy	AI independently plans, executes, and iterates within human-defined objectives and constraints.
L5	Full Autonomy	AI also identifies research objectives and coordinates sustained improvement across R&D stages without routine human direction.




Figure 14 AI4AI autonomy across R&D tasks, based on core contributor assessments. Bars show mean ratings, error bars indicate ± 1 sample standard deviation across nine contributors, and points represent individual ratings. Vertical offsets separate coincident points and have no quantitative meaning. Levels L1–L5 are encoded as 1–5; zero is a plotting baseline, not a rating category. Error bars describe inter-contributor variation, not confidence intervals. The right column gives the autonomy level the team assigned to each task, with the number of filled markers indicating the level. Assignments are consensus judgments informed by the ratings.

Autonomy is uneven across tasks (Figure 14). The team assigns each task a level by consensus: experimentation and monitoring and deployment engineering are L4; model architecture design and learning algorithm design are L2, and neither receives any rating above L3; the remaining seven tasks are L3. Contributors see substantially more scope for autonomous execution in operational tasks than in design tasks.

These are contributor judgments, not standardized measurements. Only one of the 99 ratings is L5, and no task is assigned above L4. Because the assigned levels range from L2 to L4, autonomy is better described per task than as a single maturity level for the whole workflow.

Finding 8 (AI Autonomy Is Task-Dependent): Shared context, reusable skills, and integrated tools let agents iterate with limited human intervention, but autonomy is uneven: operational tasks such as experimentation and deployment reach L4, whereas model architecture and learning algorithm design sit at L2 and never exceed L3.

7 Conclusion, Limitations, and Future Directions

7.1 Conclusion

In this report, we introduced ZGCM-1, a fully open-source 7.39B dense foundation model optimized for complex mathematical reasoning and agentic search across a 256K context. Challenging the prevailing assumption that frontier reasoning strictly necessitates hundreds of billions of parameters, ZGCM-1 is established upon a clear core thesis: *compact models can overcome inherent parametric capacity limits by coupling deliberate internal reasoning with active external tool use.*

To realize this vision on accessible academic compute, we developed an end-to-end, high-efficiency open recipe:

- **System-Architecture Co-Design:** Interleaving gated sliding-window attention with global attention at a 5:1 ratio delivers a $3.94\times$ throughput speedup and a $6.4\times$ KV-cache reduction at 256K. Paired with FP8 precision, the Muon optimizer, and TWEO outlier suppression, our pipeline achieves a $\sim 4.2\times$ pre-training time-to-loss acceleration over standard BF16/AdamW baselines.
- **Progressive Curriculum & MDP Mid-Training:** Progressively expanding context windows across 600B tokens ($16K \rightarrow 64K \rightarrow 256K$) alongside Markov Decision Process (MDP) state-action modeling injects dense per-step decision supervision into raw pre-training.
- **Calibrated Post-Training & AI-Native R&D:** Rigorous tiered data pruning and unified think/no-think co-training deliver robust reasoning and direct-response efficiency from moderate-length SFT data, complemented by a mixed RL stage that yields gains on single-domain tasks such as mathematics and code. Concurrently, an autonomous AI-native multi-agent workflow drastically compressed iteration cycles.

Extensive evaluations demonstrate that ZGCM-1-7B achieves competitive or state-of-the-art results among sub-10B models on challenging benchmarks (e.g., 75.0% on AIME 2026, 63.1% on WebWalkerQA, and 62.0% on Binary Function Search), standing toe-to-toe with models orders of magnitude larger. To foster transparent and reproducible open-science research, model weights from the pre-training, mid-training, and post-training stages, intermediate training checkpoints, data recipes, telemetry logs, and evaluation suites are made publicly available.

7.2 Limitations

Despite its competitive reasoning and agentic performance, ZGCM-1 operates under several identifiable limitations:

1. **Parametric Knowledge Bound:** While deliberate thinking and external search compensate for many factual omissions, the model’s static memory capacity remains fundamentally bounded by its 7.39B dense scale. In purely closed-book, recall-intensive tasks without retrieval support, ZGCM-1 naturally trails massive frontier systems.
2. **Instruction Adherence vs. Reasoning Verbosity:** As revealed by our fine-grained Atomic Capability Evaluation (ACE) and data mixture ablations, heavy reasoning supervision can induce verbosity and slightly impair strict, non-reasoning instruction following (e.g., IFEval and complex surface-level constraints) if not continuously calibrated.
3. **Nascent General Software and Terminal Agency:** While ZGCM-1 excels in structured agentic environments like web deep research and binary analysis, broader long-horizon repository-level engineering (e.g., SWE-bench Verified) and unstructured Linux terminal navigation (Terminal-Bench 2.0) present steep challenges where completion rates remain modest.

4. **Environment and Protocol Brittleness:** The model’s agentic execution relies heavily on strict schema alignment and stable environment feedback. Extreme observation noise, tool-calling format deviations, or external search API latency can still disrupt multi-step rollout trajectories.

7.3 Future Directions

Our findings pave the way for several high-impact directions in efficient open-source intelligence:

1. **Extension to Sparse MoE Architectures:** Scaling the hybrid SWA, Muon optimization, and MDP mid-training recipes to sparse Mixture-of-Experts (MoE) backbones will allow expanding parametric capacity and domain specialization while preserving low inference FLOPs.
2. **End-to-End Interactive Agentic RL:** Transitioning from token-level SFT and outcome-based reasoning rewards toward multi-turn, interactive reinforcement learning directly inside real-world execution sandboxes (e.g., terminal, web, and compiler environments).
3. **Autonomous Dynamic Knowledge Retrieval:** Further unifying pre-training representations with on-the-fly autonomous retrieval, enabling the model to dynamically trigger search sub-routines whenever parametric uncertainty is detected.
4. **Self-Evolving AI4AI R&D Ecosystem:** Extending the AI-native agent harness from cluster telemetry, data curation, and atomic evaluation into autonomous hypothesis formulation, automatic kernel optimization, and closed-loop synthetic environment design.

8 Contributions

Project Lead. Jiyan He.

Core Contributors (alphabetical). Guang Liang, Hao Liu, Haoxiang Guan, Jinbo Sun, Junyi Guo, Wenjun Feng, Yantai Xie, Yifei Shen.

Contributors (alphabetical). Bin Shao, Chuyang Wei, Kai Chen, Kexin Zhou, Minghang Zhu, Shuxin Zheng, Tie-Yan Liu, Taine Zhao, Wenhui Zhu, Xueyin Xu, Xiaoqing Zhang, Yatao Li, Yuxuan Ren.

References

- Joshua Ainslie, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebrón, and Sumit Sanghai. GQA: Training generalized multi-query transformer models from multi-head checkpoints. *arXiv preprint arXiv:2305.13245*, 2023. <https://arxiv.org/abs/2305.13245>.
- Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, Gabriel Martín Blázquez, Guilherme Penedo, Lewis Tunstall, Andrés Marafioti, Hynek Kydlíček, et al. SmolLM2: When smol goes big—data-centric training of a small language model, 2025. <https://arxiv.org/abs/2502.02737>.
- Allen Institute for AI. olmOCR-peS2o. Hugging Face dataset release, 2025. <https://huggingface.co/datasets/allenai/olmOCR-peS2o-0225>.
- Anthropic. Discovering cryptographic weaknesses with Claude. Blog post, July 2026. <https://www.anthropic.com/research/discovering-cryptographic-weaknesses>.
- Zhangir Azerbayev, Hailey Schoelkopf, Keiran Paster, Marco Dos Santos, Stephen McAleer, Albert Q. Jiang, Jia Deng, Stella Biderman, and Sean Welleck. Llemma: An open language model for mathematics, 2023. <https://arxiv.org/abs/2310.10631>.
- Iz Beltagy, Matthew E Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020. <https://arxiv.org/abs/2004.05150>.
- Rewon Child, Scott Gray, Alec Radford, and Ilya Sutskever. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019. <https://arxiv.org/abs/1904.10509>.
- François Chollet. On the measure of intelligence. *arXiv preprint arXiv:1911.01547*, 2019. <https://arxiv.org/abs/1911.01547>.

- Tri Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023. <https://arxiv.org/abs/2307.08691>.
- DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024. <https://arxiv.org/abs/2405.04434>.
- DeepSeek-AI. DeepSeek-V4: Towards highly efficient million-token context intelligence. *arXiv preprint arXiv:2606.19348*, 2026. <https://arxiv.org/abs/2606.19348>.
- Mostafa Dehghani, Josip Djolonga, Basil Mustafa, Piotr Padlewski, Jonathan Heek, et al. Scaling vision transformers to 22 billion parameters. *arXiv preprint arXiv:2302.05442*, 2023. <https://arxiv.org/abs/2302.05442>.
- Jingxuan Fan, Sarah Martinson, Erik Y. Wang, Kaylie Hausknecht, Jonah Brenner, Danxian Liu, Nianli Peng, Corey Wang, and Michael P. Brenner. HARDMath: A benchmark dataset for challenging problems in applied mathematics. *arXiv preprint arXiv:2410.09988*, 2024. <https://arxiv.org/abs/2410.09988>.
- Gemma Team. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025. <https://arxiv.org/abs/2503.19786>.
- GLM-5-Team. GLM-5: from vibe coding to agentic engineering. *arXiv preprint arXiv:2602.15763*, 2026. <https://arxiv.org/abs/2602.15763>.
- Harbor Framework Team. Harbor: A framework for evaluating and optimizing agents and models in container environments. Software release, 2026. <https://github.com/harbor-framework/harbor>.
- Harvard-MIT Mathematics Tournament. HMMT february tournament archives. Competition archive, 2026. <https://www.hmmt.org/www/archive/problems>. February 2025 and 2026 tournaments; accessed 2026-08-06.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, et al. OlympiadBench: A challenging benchmark for promoting AGI with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, 2024. <https://arxiv.org/abs/2402.14008>.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, et al. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021. <https://arxiv.org/abs/2009.03300>.
- Kaixuan Huang, Jiacheng Guo, Zihao Li, Xiang Ji, Jiawei Ge, Wenzhe Li, Yingqing Guo, Tianle Cai, Hui Yuan, Runzhe Wang, et al. MATH-Perturb: Benchmarking LLMs’ math reasoning abilities against hard perturbations. *arXiv preprint arXiv:2502.06453*, 2025. <https://arxiv.org/abs/2502.06453>.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, et al. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. *arXiv preprint arXiv:2403.07974*, 2024. <https://arxiv.org/abs/2403.07974>.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *The Twelfth International Conference on Learning Representations*, 2024. <https://openreview.net/forum?id=VTF8yNQM66>.
- Keller Jordan, Yuchen Jin, Vlado Boza, Jiacheng You, Franz Cesista, Laker Newhouse, and Jeremy Bernstein. Muon: An optimizer for hidden layers in neural networks. Blog post, 2024. <https://kellerjordan.github.io/posts/muon/>. Accessed 2026-09-08.
- Kimi Team. Kimi k3: Open frontier intelligence. *arXiv preprint arXiv:2607.24653*, 2026. <https://arxiv.org/abs/2607.24653>.
- Hynek Kydlíček, Guilherme Penedo, and Leandro von Werra. FinePDFs-Edu. Hugging Face dataset release, 2025. <https://huggingface.co/datasets/HuggingFaceFW/finepdfs-edu>.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, et al. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024. <https://arxiv.org/abs/2411.15124>.
- Katherine Lee, Daphne Ippolito, Andrew Nystrom, Chiyuan Zhang, Douglas Eck, Chris Callison-Burch, and Nicholas Carlini. Deduplicating training data makes language models better. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 8424–8445, 2022. <https://arxiv.org/abs/2107.06499>.
- Guang Liang, Jie Shao, Ningyuan Tang, Xinyao Liu, and Jianxin Wu. TWEO: Transformers without extreme outliers enables FP8 training and quantization for dummies. *arXiv preprint arXiv:2511.23225*, 2025. <https://arxiv.org/abs/2511.23225>.

- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, et al. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023. <https://arxiv.org/abs/2305.20050>.
- Wang Ling, Dani Yogatama, Chris Dyer, and Phil Blunsom. Program induction by rationale generation: Learning to solve and explain algebraic word problems. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics*, pages 158–167, 2017. doi: 10.18653/v1/P17-1015. <https://aclanthology.org/P17-1015/>.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems*, 2023. <https://arxiv.org/abs/2305.01210>.
- Jingyuan Liu, Jianlin Su, Xingcheng Yao, Zhejun Jiang, Guokun Lai, et al. Muon is scalable for LLM training. *arXiv preprint arXiv:2502.16982*, 2025. <https://arxiv.org/abs/2502.16982>.
- Qian Liu, Xiaosen Zheng, Niklas Muennighoff, Guangtao Zeng, Longxu Dou, Tianyu Pang, Jing Jiang, and Min Lin. RegMix: Data mixture as regression for language model pre-training. *arXiv preprint arXiv:2407.01492*, 2024. <https://arxiv.org/abs/2407.01492>.
- LLM-Core-Team Xiaomi. MiMo: Unlocking the reasoning potential of language model – from pretraining to posttraining. *arXiv preprint arXiv:2505.07608*, 2025. <https://arxiv.org/abs/2505.07608>.
- Thang Luong, Dawsen Hwang, Hoang H. Nguyen, Golnaz Ghiasi, Yuri Chervonyi, Insuk Seo, Junsu Kim, Garrett Bingham, Jonathan Lee, Swaroop Mishra, et al. Towards robust mathematical reasoning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025. <https://aclanthology.org/2025.emnlp-main.1794/>.
- Mathematical Association of America. American invitational mathematics examination. Mathematics competition, 2026. https://artofproblemsolving.com/wiki/index.php/AIME_Problems_and_Solutions. 2024, 2025, and 2026 competitions; problems and solutions archived by the Art of Problem Solving; accessed 2026-09-08.
- Mike A. Merrill, Alexander G. Shaw, Nicholas Carlini, Boxuan Li, Harsh Raj, et al. Terminal-Bench: Benchmarking agents on hard, realistic tasks in command line interfaces. *arXiv preprint arXiv:2601.11868*, 2026. <https://arxiv.org/abs/2601.11868>. Introduces Terminal-Bench 2.0, an 89-task benchmark; release announcement at <https://www.tbench.ai/news/announcement-2-0>.
- Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. GAIA: A benchmark for general AI assistants. In *International Conference on Learning Representations*, 2024. <https://arxiv.org/abs/2311.12983>.
- Paulius Micikevicius, Dusan Stolic, Neil Burgess, Marius Cornea, Pradeep Dubey, et al. FP8 formats for deep learning. *arXiv preprint arXiv:2209.05433*, 2022. <https://arxiv.org/abs/2209.05433>.
- National Security Agency. Ghidra software reverse engineering framework. Software release, 2019. <https://github.com/NationalSecurityAgency/ghidra>.
- NVIDIA Corporation. Nemotron-CC-Code-v1. Hugging Face dataset release, 2025a. <https://huggingface.co/datasets/nvidia/Nemotron-CC-Code-v1>.
- NVIDIA Corporation. Nemotron Pretraining Code v1 and v2. Hugging Face dataset releases, 2025b. <https://huggingface.co/datasets/nvidia/Nemotron-Pretraining-Code-v2>. See also Nemotron-Pretraining-Code-v1.
- NVIDIA Corporation. Nemotron-Pretraining-Specialized-v1. Hugging Face dataset release, 2025c. <https://huggingface.co/datasets/nvidia/Nemotron-Pretraining-Specialized-v1>.
- NVIDIA Corporation. Megatron-LM and Megatron Core: Gpu-optimized library for training transformer models at scale. Software repository and documentation, 2026. <https://github.com/NVIDIA/Megatron-LM>. Accessed 2026-08-06.
- OpenAI. Introducing SWE-bench Verified. OpenAI blog, 2024. <https://openai.com/index/introducing-swe-bench-verified/>. Accessed 2026-09-08.
- OpenAI. GPT-5.6: Frontier intelligence that scales with your ambition. Model release, July 2026. <https://openai.com/index/gpt-5-6/>.
- Jake Poznanski, Aman Rangapur, Jon Borchardt, Jason Dunkelberger, Regan Huff, Daniel Lin, Christopher Wilhelm, Kyle Lo, and Luca Soldaini. olmOCR: Unlocking trillions of tokens in PDFs with vision language models, 2025. <https://arxiv.org/abs/2502.18443>.

- Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi. Generalizing verifiable instruction following. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2025. <https://arxiv.org/abs/2507.02833>.
- Zihan Qiu, Zekun Wang, Bo Zheng, Zeyu Huang, Kaiyue Wen, Songlin Yang, Rui Men, Le Yu, Fei Huang, Suozhi Huang, Dayiheng Liu, Jingren Zhou, and Junyang Lin. Gated attention for large language models: Non-linearity, sparsity, and attention-sink-free, 2025. <https://arxiv.org/abs/2505.06708>.
- Qwen Team. Qwen3.8-Max: A new bar for coding and cowork. Blog post, August 2026. <https://qwen.ai/blog?id=qwen3.8>.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, et al. GPQA: A graduate-level google-proof Q&A benchmark. In *First Conference on Language Modeling*, 2024. <https://arxiv.org/abs/2311.12022>.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, et al. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. <https://arxiv.org/abs/2402.03300>.
- Noam Shazeer. GLU variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020. <https://arxiv.org/abs/2002.05202>.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. <https://arxiv.org/abs/2104.09864>.
- Liangcai Su, Zhen Zhang, Guangyu Li, Zhuo Chen, Chenxi Wang, Maojia Song, Xinyu Wang, Kuan Li, Jialong Wu, Xuanzhong Chen, et al. Scaling agents via continual pre-training. *arXiv preprint arXiv:2509.13310*, 2025. <https://arxiv.org/abs/2509.13310>.
- Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V. Le, Ed H. Chi, Denny Zhou, and Jason Wei. Challenging BIG-Bench tasks and whether chain-of-thought can solve them. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13003–13051, 2023. doi: 10.18653/v1/2023.findings-acl.824. <https://arxiv.org/abs/2210.09261>.
- SWE-agent. mini-SWE-agent. Software repository, version v2, 2026. <https://github.com/SWE-agent/mini-swe-agent>. Accessed 2026-09-08.
- Team Olmo. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025. <https://arxiv.org/abs/2512.13961>.
- Together Computer. RedPajama-Data-1T: An open source recipe to reproduce the LLaMA training dataset. Hugging Face dataset release, 2023. <https://huggingface.co/datasets/togethercomputer/RedPajama-Data-1T>. ArXiv slice, 28B tokens.
- vLLM Team. vLLM: A high-throughput and memory-efficient inference and serving engine for LLMs. Software repository, 2023. <https://github.com/vllm-project/vllm>. See also Kwon et al., SOSP 2023, arXiv:2309.06180.
- Yudong Wang, Zixuan Fu, Jie Cai, Peijun Tang, Hongya Lyu, Yewei Fang, Zhi Zheng, Jie Zhou, Guoyang Zeng, Chaojun Xiao, Xu Han, and Zhiyuan Liu. Ultra-FineWeb: Efficient data filtering and verification for high-quality LLM training data, 2025. <https://arxiv.org/abs/2505.05427>.
- Jason Wei, Maarten Bosma, Vincent Y. Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. Finetuned language models are zero-shot learners. In *International Conference on Learning Representations*, 2022. <https://arxiv.org/abs/2109.01652>.
- Jason Wei, Zhiqing Sun, Spencer Papay, Scott McKinney, Jeffrey Han, Isa Fulford, Hyung Won Chung, Alex Tachard Passos, William Fedus, and Amelia Glaese. BrowseComp: A simple yet challenging benchmark for browsing agents. *arXiv preprint arXiv:2504.12516*, 2025. <https://arxiv.org/abs/2504.12516>.
- Jialong Wu, Baixuan Li, Runnan Fang, Wenbiao Yin, Liwen Zhang, Zhengwei Tao, Dingchu Zhang, Zekun Xi, Gang Fu, Yong Jiang, Pengjun Xie, Fei Huang, and Jingren Zhou. WebDancer: Towards autonomous information seeking agency. *arXiv preprint arXiv:2505.22648*, 2025a. <https://arxiv.org/abs/2505.22648>.
- Jialong Wu, Wenbiao Yin, Yong Jiang, Zhenglin Wang, Zekun Xi, et al. WebWalker: Benchmarking LLMs in web traversal. *arXiv preprint arXiv:2501.07572*, 2025b. <https://arxiv.org/abs/2501.07572>.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025a. <https://arxiv.org/abs/2505.09388>.

- An Yang, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2.5-1M technical report. *arXiv preprint arXiv:2501.15383*, 2025b. <https://arxiv.org/abs/2501.15383>.
- John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik R. Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. <https://arxiv.org/abs/2405.15793>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, et al. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023. <https://arxiv.org/abs/2210.03629>.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025. <https://arxiv.org/abs/2503.14476>.
- Biao Zhang and Rico Sennrich. Root mean square layer normalization. *Advances in Neural Information Processing Systems*, 32, 2019. <https://arxiv.org/abs/1910.07467>.
- Yang Zhang, Amr Mohamed, Hadi Abdine, Guokan Shang, and Michalis Vazirgiannis. Beyond random sampling: Efficient language model pretraining via curriculum learning. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5776–5794. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.eacl-long.271. <https://aclanthology.org/2026.eacl-long.271/>.
- Kunhao Zheng, Jesse Michael Han, and Stanislas Polu. miniF2F: A cross-system benchmark for formal olympiad-level mathematics. In *International Conference on Learning Representations*, 2022. <https://arxiv.org/abs/2109.00110>.
- Wanjun Zhong, Ruixiang Cui, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang, Amin Saied, Weizhu Chen, and Nan Duan. AGIEval: A human-centric benchmark for evaluating foundation models. *arXiv preprint arXiv:2304.06364*, 2023. <https://arxiv.org/abs/2304.06364>.
- Chuyue Zhou, Hongya Lyu, Xinle Lin, Hengyu Zhao, Junshao Guo, Xueren Zhang, Shuaikang Xue, Qiang Ma, Jie Zhou, Yudong Wang, and Zhiyuan Liu. UltraData-Math. Hugging Face dataset release, 2026. <https://huggingface.co/datasets/openbmb/UltraData-Math>.
- Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, et al. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023. <https://arxiv.org/abs/2311.07911>.

Contents

1	Introduction	2
2	Architecture	3
2.1	Model Overview	3
2.2	Hybrid-Attention Experiments	4
2.3	KV Cache Analysis	5
3	Pre-Training	5
3.1	General Pre-Training	5
3.1.1	Data Mixture	5
3.1.2	Curriculum Pretraining	6
3.1.3	Hyperparameters	7
3.2	Mid-Training	7
3.2.1	Data Mixture	7
3.2.2	Reasoning Data	9
3.2.3	Instruction Data	9
3.2.4	Agentic Data	9
3.2.5	Hyperparameters	10
3.3	Monitor During Training	10
4	Post-Training	11
4.1	Supervised Fine-Tuning	11
4.1.1	General Data	11
4.1.2	Agentic Data	13
4.1.3	Training Objective and Hyperparameters	14
4.2	Mixed Reinforcement Learning	15
4.2.1	RL Data Curation	15
4.2.2	Reward Design	15
4.2.3	Hyperparameters	15
5	Evaluation	15
5.1	Setup	15
5.2	7B Scale Model Comparison	16
5.3	Agentic Evaluation	16

6	AI-Native Research and Development	18
6.1	AI Assistance Across the R&D Lifecycle	18
6.2	AI4AI Autonomy Across the R&D Lifecycle	19
7	Conclusion, Limitations, and Future Directions	21
7.1	Conclusion	21
7.2	Limitations	21
7.3	Future Directions	22
8	Contributions	22
A	Pre-Training Details	30
A.1	Training Configurations	30
A.2	Capability Dynamics during General Pre-Training	30
A.3	Capability Development during Mid-Training	32
A.4	Curriculum Pretraining	32
A.4.1	Qualitative Examples: Why Code and Mathematics Are Interleaved Separately	33
A.5	Mid-Training Mixtures	34
A.5.1	Reasoning and Instruction Data	34
A.5.2	Agentic Data	34
B	Post-Training Implementation Details	35
B.1	General SFT Data Selection	35
B.1.1	Evidence for Quality-First Selection	35
B.2	Agentic SFT Data	35
B.2.1	Deep-Research Trajectories	35
B.2.2	Software-Engineering Trajectories	36
B.2.3	Terminal Trajectories	37
B.3	Packing Validation	37
B.4	Thinking-to-Direct Transfer	37
C	Agentic Evaluation Protocol	38
C.1	Web-Environment Deep-Research Harness	38
C.2	Binary Function Search	38
C.3	SWE-bench Verified Mini50 Harness	40
C.4	Terminal-Bench 2.0 Harness	41

D Atomic Capability Evaluation Framework	41
D.1 Design Principles	41
D.2 Taxonomy and Scoring	41
D.3 AI-Assisted Benchmark Construction	42
D.4 Evaluation Results	42
D.5 Quality and Scope	43
E AI Autonomy Rubric and Contributor Assessments	43
E.1 Assessment Scope and Interpretation	44
E.2 Task-Specific Autonomy Criteria	44
E.2.1 Data Cleaning	44
E.2.2 Data Acquisition	44
E.2.3 Synthetic Data Generation	45
E.2.4 Model Architecture Design	45
E.2.5 Learning Algorithm Design	45
E.2.6 Experimentation and Monitoring	46
E.2.7 Operator Design	46
E.2.8 Training and Inference Framework Design	46
E.2.9 Resource Platform Management	46
E.2.10 Model Evaluation	47
E.2.11 Deployment Engineering	47
E.3 Individual Contributor Ratings	47
F Evaluation Reproducibility and Provenance	48

Appendix

A Pre-Training Details

A.1 Training Configurations

Table 6 lists the distributed-training configurations for general pre-training (16K context) and long-context qualification (64K and 256K). All runs use H100 GPUs with FP8 computation. We evaluate two routes to 256K: a direct route that trains 30B tokens at 256K, and a staged route that first trains 10B tokens at 64K then 20B tokens at 256K. The staged route reaches a slightly lower final loss (1.16 vs. 1.19) at comparable throughput.

Table 6 Distributed-training configurations on H100 GPUs. TP = tensor parallelism, PP = pipeline parallelism, CP = context parallelism, DP = data parallelism, MBS = micro-batch size, GBS = global batch size.

Stage	Context	TP/PP/CP/DP	MBS/GBS	Recompute	LR	RoPE base	Notes
General pre-training	16K	2/1/1/96	1/768	None	2×10^{-4}	5M	~60% BF16-eq. MFU (≈ 585 TFLOP/s/GPU)
Staged (step 1)	64K	2/1/4/24	1/192	Selective	1×10^{-4}	5M	10B tokens
Direct 256K	256K	2/2/4/12	1/48	Full	1×10^{-4}	10M	30B tokens; 203 TFLOP/s/GPU
Staged (step 2)	256K	2/2/4/12	1/48	Full	1×10^{-4}	10M	20B tokens; 204 TFLOP/s/GPU

A.2 Capability Dynamics during General Pre-Training

We directly evaluate 54 saved base checkpoints, from 0.04T to 4.19T cumulative General Pre-Training tokens processed. Each checkpoint is evaluated in its base-model form and generates responses without supervised fine-tuning, instruction tuning, or other post-training adaptation. The benchmark harness and evaluation settings remain fixed across checkpoints, so the curves reflect capability changes in the pre-trained model itself. Token positions are read from the training logs as consumed samples times sequence length, which keeps them exact across the global-batch-size ramp. The resulting trajectories are shown in Figure 15.

Evaluation uses fixed subsets rather than full test sets: 1,425 items for MMLU (25 per subject), 1,004 for HellaSwag, and between 50 and 200 items for the remaining benchmarks. GSM-Symbolic, Minerva Math, and OpenBookQA use 50 items each, so a single item shifts their score by two points. Point-to-point variation on these three benchmarks reflects sampling noise as much as capability change, and only their overall trend is informative.

Table 7 summarizes four five-checkpoint windows. The first and last columns average the first and last five checkpoints; the two middle columns average five checkpoints centered near 1.0T and 2.5T tokens. Scores are reported as percentages.

Every benchmark improves from the first window to the last, but the rate at which they improve diverges. Most of the gain arrives before 2.5T tokens: MMLU covers 92% of its total improvement by the 2.5T window, ARC-Challenge 88%, and WinoGrande 84%. The final 1.7T tokens add 2.58 points on MMLU, 1.97 on GSM8K, and 2.84 on WinoGrande. Code and the harder mathematics benchmarks keep improving over the same interval, with HumanEval adding 7.08 points and GSM-Symbolic 14.80, though the latter is measured on 50 items and carries wide uncertainty. These diverging dynamics suggest that broad pre-training alone provides increasingly uneven marginal returns and motivate the capability-oriented mixture used in Mid-Training.




Figure 15 Capability trends during General Pre-Training. Each point is obtained by directly evaluating the corresponding base checkpoint without supervised fine-tuning or other post-training adaptation. The x-axis reports cumulative tokens processed, truncated at 4.19T. The dotted line at 1.0T marks the end of the curriculum-ordered stage. Light lines show individual evaluations, and dark lines show five-point moving averages.

Table 7 Five-checkpoint window averages during General Pre-Training. Higher is better.

Benchmark	Capability	Early (0.04T–0.07T)	Around 1.0T (0.91T–1.16T)	Around 2.5T (2.29T–2.79T)	Final (3.67T–4.18T)
MMLU	Knowledge	26.37	48.32	56.94	59.52
OpenBookQA	QA	24.80	60.80	72.40	78.80
CSQA	QA	19.02	57.70	64.26	67.87
HumanEval	Code	13.66	28.29	36.46	43.54
GSM8K	Math	11.52	38.64	57.12	59.09
Minerva Math	Math	6.80	20.40	32.80	38.40
GSM-Symbolic	Math	1.20	25.60	49.20	64.00
ARC-Challenge	Reasoning	24.96	56.07	69.74	75.90
HellaSwag	Reasoning	24.56	42.75	52.69	57.15
PIQA	Reasoning	53.80	64.78	69.78	74.89
WinoGrande	Reasoning	50.71	65.98	66.14	68.98

A.3 Capability Development during Mid-Training

We use a short-SFT probe to track the downstream capability of successive training milestones. Each checkpoint is rapidly adapted with the same carefully curated SFT corpus and then evaluated on a fixed benchmark suite. The probe measures the capability elicited from each initialization after lightweight supervised adaptation. It is therefore distinct from the direct base-checkpoint evaluation in [Appendix A.2](#).




Figure 16 Capability trajectories across General Pre-Training and Mid-Training milestones after rapid fine-tuning on the same SFT corpus. The General Pre-Training probe and the 15%/30%/70%/77% Mid-Training probes use two SFT epochs; the final (100%) Mid-Training probe uses three. The curves show the downstream capability exposed by short SFT rather than direct base-model performance.

The probe shows a large shift after entering Mid-Training. From the end of General Pre-Training to the end of Mid-Training, MATH-500 rises from 35.83% to 74.12%, HumanEval+ from 43.29% to 73.78%, IFEval Strict from 29.57% to 72.64%, and MMLU from 44.17% to 72.63%. Individual benchmarks are not monotonic at every checkpoint.

Because the final Mid-Training probe uses three SFT epochs while the preceding probes use two, these trajectories are diagnostic evidence of improving downstream potential rather than a strictly matched causal ablation. Together with the direct base-model trends in [Appendix A.2](#), they support the use of capability-oriented Mid-Training before full post-training.

A.4 Curriculum Pretraining

We compare two 7B runs trained for 10B tokens with identical data mixtures. The random baseline samples documents in random order. For the curriculum run, we first exclude documents with extreme lexical-complexity scores as outliers; these cases are typically corrupted or garbled text rather than meaningful difficult examples. We then order the remaining natural-language documents from lower to higher lexical complexity. Code and mathematics do not use lexical complexity as a difficulty proxy; they are shuffled independently and interleaved at their target proportions.

Table 8 Bits per byte (BPB) after 10B training tokens. Lower is better.

Benchmark	Random	Curriculum
Coding	1.99	0.81
Math	0.97	0.94
MMLU	1.03	1.12
ARC	1.05	1.09
HellaSwag	1.05	1.08

Curriculum ordering improves coding BPB (1.99 $\rightarrow$ 0.81) and math BPB (0.97 $\rightarrow$ 0.94). General benchmarks

(MMLU, ARC, HellaSwag) show a modest increase of 0.03–0.09 BPB, suggesting that the curriculum accelerates domain-specific learning at a small cost to general knowledge acquisition at this token budget.

A.4.1 Qualitative Examples: Why Code and Mathematics Are Interleaved Separately

The examples below illustrate a limitation of using lexical complexity as a single ordering signal. Percentile positions are recomputed within the corresponding source shard and are reported only as diagnostic positions; they are not intrinsic difficulty labels or production metadata.

Code. The first example is a WFC-based grid-construction program. It appears near the low end of the ordering (first 0.01% of the code shard), despite containing substantial program structure and environment-specific logic.

Code example: low end of the ordering

```
import wfc_implementation as wfci
import mapUtils, interfaceUtils
buildArea = interfaceUtils.requestBuildArea()
tiles = np.array([[0,0,0,0], [0,1,0,0], [1,1,0,1]])
```

Diagnostic position: first 3 of 25,000 documents.

By contrast, the following two-line NumPy snippet appears near the high end of the ordering (last 0.01%), although it performs only a simple file-loading operation.

Code example: high end of the ordering

```
import numpy as np
count_data = np.loadtxt("../BookSource/Chapter1_Introduction/data/txtdata.csv")
```

Diagnostic position: last 3 of 25,000 documents.

These two examples show that surface lexical complexity can be dominated by imports, boilerplate, and token-level variation rather than by the algorithmic content of a program.

Mathematics. The same issue occurs in mathematical text. A long combinatorial programming problem, “Manhattan Wiring,” appears near the low end of the ordering (first 0.11% of the mathematics shard), even though it requires non-intersecting path construction under obstacles and a minimum-length objective.

Mathematics example: low end of the ordering

Connect two pairs of marked cells with non-intersecting horizontal or vertical lines around obstacles, while minimizing the total length of the two lines.

Diagnostic position: first 63 of 55,743 documents.

At the opposite end, a one-line statement of the associative law appears near the high end of the ordering (last 0.01%), despite requiring little mathematical reasoning.

Mathematics example: high end of the ordering

This shows the associative property of addition: $(a + b) + c = a + (b + c)$.

Diagnostic position: last 3 of 55,743 documents.

The mathematics examples reinforce the same design choice: lexical complexity is useful for efficiently ordering general-language data, but it is not a reliable proxy for the reasoning difficulty of code or mathematical content. We therefore interleave these two domains separately throughout curriculum pretraining.

A.5 Mid-Training Mixtures

Mid-training uses continued pre-training with a full-sequence causal objective (not chat-template supervision). Instruction records are serialized as raw context, and agent actions and environment observations remain next-token targets.

We apply pool-wide deduplication and mathematics error filtering to produce a candidate pool of approximately 5.3B records (2.86T tokens). The pool is then sampled into three cumulative context mixtures:

Table 9 Mid-training context mixtures. Each longer-context stage retains shorter-context data.

Stage	Sampled tokens	Length composition
16K	180B	180B at $\leq 16K$
64K	240B	180.89B at $\leq 16K$ + 59.11B at 16K–64K
256K	180.51B	127.81B at $\leq 16K$ + 21.72B at 16K–64K + 30.98B above 64K

Code and mathematics each occupy approximately 20% across all three stages. The 256K stage increases knowledge from 10.5% to 14.2% and agentic data from 1.5% to 3.3%, while reducing pre-training replay from 13% to 9%.

A.5.1 Reasoning and Instruction Data

Reasoning sources are normalized to a common schema and classified into accepted, rewrite, pending, holdout, or rejected categories based on self-containment, answer evidence, and reasoning quality. The reasoning pool contains approximately 35M examples (38B tokens). Of these, 34.7M are retained directly via rule-based cleaning; 201K examples with valuable questions but low-quality traces are reconstructed by teacher models and re-admitted after validation.

Instruction data is collected from multiple sources (including Tulu- and FLAN-derived data) (Lambert et al., 2024; Wei et al., 2022) and converted into pre-training-compatible context. Selected single-turn examples are rewritten as multi-turn discussions to provide supervision for sequential reasoning.

A.5.2 Agentic Data

Agentic mid-training data combines general interaction traces with software-engineering trajectories.

For general interaction traces, we retain high-quality complete interactions that cover long-horizon task context, tool calls, and environment feedback. We additionally reformulate retained traces as Markov decision process-style single-step decision examples. At each decision point, the task description, interaction history, and current environment feedback form the state, and the model predicts an appropriate next action conditioned on that state. This shifts the learning target from reproducing complete state-transition sequences to state-conditioned action selection, providing denser per-step supervision alongside the full-trajectory signal.

For software-engineering trajectories, we apply different filtering strategies depending on whether a real execution environment is available. Execution-free trajectories are screened using structural, patch, shell/edit, log, and behavioral evidence rather than self-reported runtime outcomes. Execution-grounded trajectories are jointly scored and filtered according to actual execution results, feedback from the final interaction turn, and the number of interaction turns. Retained trajectories preserve task context, tool calls, observations, code changes, and final responses. After deduplication, the final software-engineering trajectory corpus contains approximately 129K records (5.9B tokens). These trajectories are trained with full-sequence causal language modeling, where both agent actions and environment observations are next-token targets.

B Post-Training Implementation Details

B.1 General SFT Data Selection

General SFT data is assembled from open-source datasets and internally distilled examples spanning instruction following, knowledge, code, mathematics, science, dialogue, and reasoning. The shared quality-selection and decontamination procedures are described in [Section 4.1](#); this appendix records domain-specific selection and implementation details.

We compare candidate domain mixtures across instruction following, reasoning, and other capabilities. Excessive long-reasoning shares reduce instruction following, so we iteratively adjust long-reasoning, direct-response, and instruction-following proportions before selecting the final composition.

Within the general SFT data, the reasoning component retains 1,307,175 of 2,268,178 candidate examples after structural validation, answer-quality review, and reasoning-value filtering. Accepted examples must contain coherent, complete, and substantively useful reasoning in addition to a reliable final answer. Safety and multi-turn examples pass task-appropriate structural checks, and reasoning content remains separate from visible final answers throughout processing.

B.1.1 Evidence for Quality-First Selection

To isolate the effect of data quality, we start from the same base checkpoint and quickly fine-tune it into a conversational model using three versions of a common SFT pool. The first version receives only format normalization and basic structural checks, without model-based quality filtering. The second removes examples assigned to review or deletion while retaining broad capability coverage. The third keeps only examples that pass the strictest quality, style, and capability checks. These counts refer to this controlled ablation pool rather than to the final SFT corpus.

Table 10 Ablation of SFT data filtering. All variants use the same base checkpoint, rapid SFT procedure, and evaluation settings. The final column is the mean over the six shared benchmarks.

Data processing	Samples	IFEval	MATH-500	BBH	MMLU	HumanEval+	MBPP+	Mean (6)
Minimal processing	≈2.08M	72.64	74.12	58.49	72.63	73.78	55.03	67.78
Broader filtering	≈1.833M	71.20	77.63	65.08	70.60	65.00	58.36	67.98
Quality-focused filtering	≈1.145M	71.94	75.43	68.59	71.50	67.56	57.94	68.83

The quality-focused version uses 44.9% fewer examples than the minimally processed version and 37.5% fewer than the broadly filtered version. Despite its smaller size, its mean over the six shared benchmarks rises from 67.78 and 67.98 to 68.83, respectively. Individual tasks do not all move in the same direction, but the aggregate result shows that high-quality data is more useful than simply increasing data volume.

B.2 Agentic SFT Data

We compare two schedules for introducing agentic supervision. The sequential schedule first trains on general-capability data and then continues on agentic data. The joint schedule interleaves general and agentic examples throughout SFT. Candidate-run evaluations favor the joint schedule, which is used for the final training run.

B.2.1 Deep-Research Trajectories

Interaction schema. The Deep Research system instruction is:

You are a deep search assistant. Your primary role is to perform rigorous, multi-step, multi-source investigations on any topic, covering both broad open-domain questions and highly specialized academic inquiries. For each user request, you must actively seek out and cross-check information from credible and diverse sources, then integrate the findings into a response that is comprehensive, accurate, well-structured, and objective. When you have gathered sufficient information and are ready to provide the definitive response, enclose the entire final answer in <answer></answer> tags.

The interaction exposes two tools:

Table 11 Tool schema for Deep Research trajectories.

Tool	Arguments	Return
search	query : nonempty string array; topn : positive integer	Ranked results containing titles, URLs, and snippets for each query.
visit	url : nonempty URL array; goal : natural-language information goal	Goal-directed information extracted from the requested pages.

Each example follows the message sequence `system` $\rightarrow$ `user` $\rightarrow$ `assistant` $\rightarrow$ `tool`, with assistant-tool exchanges repeated as needed before the final assistant response. Tools are declared in the top-level `tools` field; assistant actions contain a tool name and JSON arguments in `tool_calls`; and each environment return is a separate `role=tool` message paired with the preceding call. Search observations use a ranked result list, while visit observations provide information relevant to the stated goal. The final response is enclosed in `<answer>...</answer>`.

Think examples include assistant reasoning before actions and the final answer. Their aligned no-think views remove the reasoning spans while preserving the user request, tool calls, observations, and answer.

B.2.2 Software-Engineering Trajectories

Execution-grounded trajectory construction and filtering. The execution-grounded source pool contains 34,269 trajectories. We require valid nonempty tool calls, interaction with external tools, tool observations, and a complete finish action. We remove trajectories containing private paths, secret-like strings, empty completion content, or a finish message that still describes work in progress. This first pass retains 30,914 trajectories and rejects 3,355: 3,067 for private paths, 254 for secret patterns, and 34 for in-progress completion language. All 30,914 retained trajectories pass the GLM-5.1 tokenizer and supervision-boundary preflight without tokenizer errors or empty targets. Restricting the view to 8K–64K tokens removes 900 overlength trajectories and yields 30,014 execution-grounded examples.

Execution-free trajectory construction, filtering, and scoring. The execution-free source pool contains 318,115 trajectories. The initial structural and safety gate retains 281,237 candidates and rejects 36,878: 35,259 for private paths, 1,361 for secret patterns, and 258 for incomplete finish language. GLM-5.1 tokenization produces 280,852 candidates in the 8K–64K range. Structural validity alone is insufficient for execution-free data, so we construct a source-specific quality rubric through model-assisted rule iteration. Hard admission criteria require a patch, both shell and edit actions, a unique nonempty finish action, reasonable interaction and target lengths, and no unsupported claim that an unexecuted test has passed. This gate rejects 3,554 trajectories and leaves 277,298 scored candidates.

The scoring function assigns higher scores to trajectories with coherent patch evidence, complementary shell and editing behavior, balanced tool and reasoning actions, dense file-path and code-symbol evidence, sufficient assistant target tokens, informative file and terminal observations, and limited observation clipping. It assigns lower scores to trajectories with excessive turns or reasoning actions, heavy clipping, uncertain completion language, or runtime-success claims unsupported by an execution environment. We select the top 30,000 trajectories while capping each repository at 800 examples to limit source concentration. The selected subset has a minimum quality score of 97, a median of 97, and a 95th percentile of 100. An independent audit confirms that all selected examples are marked execution-free and contain no illegal tools or empty finish actions.

Structured-tool projection and validation. Quality selection is performed on traceable native records before training-format conversion. We then project the 30,014 execution-grounded and 30,000 execution-free trajectories to the GLM-5.1 top-level `tools` schema. External `execute_bash` and `str_replace_editor` actions become structured assistant `tool_calls`; native `think` actions become `reasoning_content`; and `finish` actions become the final assistant response. Internal acknowledgements such as “Your thought has

been logged.” are removed, consecutive assistant messages are merged, valid observations remain separate tool messages, and a trailing unpaired tool observation is dropped. The resulting 60,014 software-engineering candidates pass enhanced structural checks for JSON validity, top-level tool definitions, tool-call-observation pairing, duplicate identifiers, and complete assistant termination.

B.2.3 Terminal Trajectories

Environment and schema. We collect terminal-agent trajectories from terminal task environments through two independent collection routes, covering Claude Opus 4.6 Thinking and GLM-5.2. Deduplication is performed independently within each collection route using the task-environment identity. Each task runs in an isolated Docker environment with a persistent shell and one structured `bash` tool. A rollout permits at most 64 interaction steps and 4,096 output tokens per assistant turn. Commands have a 180-second timeout, and each terminal observation is capped at 12,000 characters.

The agent terminates through a predefined submission command, after which `tests/test.sh` is executed in the same environment to obtain a verifier outcome independent of the model’s own completion claim. The full trajectory, run summary, and provenance manifest are retained so that the task, collection route, execution, interaction trace, and verifier outcome remain traceable.

Construction, filtering, and source normalization. Raw rollouts contain repeated task runs, malformed tool protocols, missing containers or images, setup output, harness-generated format-repair turns, duplicated visible reasoning, and trailing unpaired observations. Within each collection route, we retain at most two valid trajectories for a task environment, preferring verifier-successful and then more complete runs. We reject repeated protocol failures and infrastructure failures because they do not represent valid task interactions.

Every assistant turn contains `content`, while native `reasoning_content` is available only for a subset of turns. Claude Opus 4.6 Thinking does not expose `reasoning_content` on every turn, whereas GLM-5.2 does. We normalize the two fields across sources without duplicating supervision. Setup output, harness `Format error` repairs, and infrastructure-failure text are removed from model-visible content. A terminal observation that would otherwise end the conversation is cropped so that the example closes with a complete assistant turn. Structurally valid unsuccessful trajectories retain their outcome labels, enabling both full and success-only materializations.

Materialized training views. For a later controlled Agent-SFT experiment, we materialize a paired-source view from the Claude Opus 4.6 Thinking and GLM-5.2 collection routes. The Claude source retains 11,653 of 11,944 raw runs, including 8,488 verifier-successful and 3,165 unsuccessful trajectories. The GLM-5.2 source retains 10,656 of 11,351 pointer-manifest records, including 7,260 successful and 3,396 unsuccessful trajectories.

B.3 Packing Validation

Before packing, a tools-aware chat-template and loss-mask preflight verifies tool rendering, assistant supervision boundaries, action-observation closure, sequence length, and nonempty targets. Accepted, overlength, and rendering-error records are routed to separate outputs rather than being silently truncated or discarded.

B.4 Thinking-to-Direct Transfer

A historical checkpoint comparison examines whether reasoning-oriented supervision can improve direct-response capability. A baseline checkpoint trained with direct-response data was subsequently retrained with a mixture of direct-response examples and explicit reasoning trajectories. Both checkpoints were evaluated in no-think mode, so the comparison measures answer quality without exposing an intermediate reasoning trace.

Table 12 No-think performance before and after mixed think/no-think SFT. Both columns use matching benchmark versions and metrics; rows are sorted by score change.

Benchmark	Before mixed SFT	After mixed SFT	Change
AIME 2025	10.00	43.33	+33.33
AIME 2024	6.67	33.33	+26.66
MBPP+	55.03	75.93	+20.90
BBH	58.49	69.48	+10.99
LiveCodeBench v3	28.00	34.31	+6.31
IFEval	72.64	78.00	+5.36
MMLU	72.63	70.03	-2.60
HumanEval+	73.78	66.46	-7.32

The largest gains occur in mathematical reasoning and code generation. AIME 2025 increases from 3/30 to 13/30, AIME 2024 increases from 2/30 to 10/30, and MBPP+ improves by 20.90 points. BBH, LiveCodeBench v3, and IFEval also improve. These results are consistent with reasoning supervision transferring useful problem-solving behavior to direct-response inference even when the intermediate reasoning is not visible.

The transfer is task-dependent rather than universal: MMLU decreases by 2.60 points and HumanEval+ by 7.32 points.

This experiment is an observational checkpoint comparison, not a controlled data ablation. The mixed-data run may differ in total training tokens, data composition, checkpoint history, and optimization schedule.

C Agentic Evaluation Protocol

This appendix documents the evaluation harnesses used for the web-environment results in Table 3, the Binary Function Search results in Table 4, and our internal software-engineering and terminal-agent diagnostics.

C.1 Web-Environment Deep-Research Harness

The deep-research evaluation uses a thinking-enabled ReAct policy with at most 64 search-and-read steps. Search queries are issued through Serper with Jina Search as a fallback, and pages are extracted through Jina Reader. Retrieved material is summarized by ZGCM-1 operating in direct-response (no-thinking) mode, while final answers are judged by Qwen3-30B-A3B-Instruct-2507. If the interaction reaches the step budget, the harness requests a final answer from the evidence accumulated so far.

Results under this protocol are reported in Table 3. Their interpretation depends on search date and availability, provider behavior, page extraction, summarization, judge settings, and retry policy.

C.2 Binary Function Search

Task Definition. Binary Function Search evaluates whether an agent can locate a target function in a stripped ELF binary from a natural-language description. Stripping removes most symbolic clues, while compiler optimization may transform the source-level structure. Moreover, several functions may share the same strings, constants, or callees. The agent therefore must distinguish functions by combining semantic and structural evidence.

Input and Output. Each task specifies the target binary, its architecture, the address-space convention, and a behavioral description of the target function. A description may include several kinds of semantic evidence:

- the function’s main behavior or transformation;
- characteristic control-flow decisions or processing stages;
- calls to other routines or interactions with subsystems;
- distinctive constants, strings, flags, or data formats;

- error detection, recovery, and reporting behavior; and
- externally visible outputs or state changes.

In most cases, the description covers only a subset of the evidence categories listed above. The following example presents input of an actual task:

```
Binary path: /chroot/build/binutils/nm
Architecture: x86_64
Address space: elf_va
```

Function description:

```
Formats one symbol record for output. It selects between
size-based and regular modes, optionally prints source
locations, resolves relocation information, classifies the
symbol, and dispatches the result to the active output-style
printer.
```

The required output is the exact function-entry ELF virtual address, together with an optional confidence level and a short rationale:

```
{
  "entry_va": "0x...",
  "confidence": "low|medium|high",
  "notes": "optional short rationale"
}
```

The submitted address must identify the beginning of the target function. File offsets, runtime addresses, and interior instructions are not valid outputs. The confidence and notes fields do not affect correctness.

Harness and Tools. The Harness provides a Ghidra-backed ([National Security Agency, 2019](#)) environment for Binary Function Search. It exposes `gcmd`, a structured interface to the results of Ghidra preprocessing. The lower panel of [Figure 12](#) summarizes its six actions:

- `functions` returns a compact, filterable list of identified functions;
- `function` returns structured metadata for a specified function;
- `decompile` returns C-like pseudocode for a specified function;
- `search` searches strings, disassembly, decompiled text, or byte sequences;
- `refs` finds references to a function, string, data object, or address; and
- `listing` returns the assembly listing around a code location.

The Harness manages the Ghidra session, normalizes addresses to the ELF virtual-address convention, and verifies that submitted addresses correspond to function boundaries. It also manages large tool outputs, detects repeated identical calls, and records the interaction for auditing.

Agentic Workflow. [Figure 12](#) summarizes the workflow, which consists of four stages.

I. Task setup. Before the first model turn, the Harness imports the stripped binary into Ghidra to recover function boundaries, control-flow structures, references, strings, assembly listings, and decompiled representations. The agent can access this information through the structured tool interface. The Harness also provides a summary of the function inventory, such as the number and size distribution of analyzed functions, without selecting or ranking candidates.

II. Candidate exploration. The agent searches the set of functions identified by Ghidra using information from the task description. It enumerates functions, applies structural filters, and searches strings, disassembly, decompiled text, or byte sequences. It then follows references from relevant search results to the functions that use them. This process reduces the function set to a shortlist of plausible candidates.

III. Evidence refinement. For each shortlisted candidate, the agent examines function metadata, decompiled code, references, callers, callees, and assembly listings. It compares the collected information with the evidence categories present in the task description, including behavior, control flow, calls, constants, errors, and outputs. The agent also examines caller-callee relationships to determine which function directly owns the described behavior. If the collected evidence is insufficient, the workflow returns to candidate exploration. The agent performs additional searches, follows new references, or updates the candidate shortlist. The two stages repeat until the evidence is sufficient for a final decision.

IV. Commit and validation. Once the agent has collected sufficient evidence and is confident in its selection, it submits the entry address of the candidate function. The Harness checks that the address corresponds to a function boundary identified during Ghidra preprocessing, normalizes it to the ELF virtual-address convention, and compares it with the hidden oracle.

During candidate exploration and evidence refinement, each agent turn produces at most one structured `gcmd` call based on the task and previous observations. The Harness executes the call and returns its result to the agent, forming an iterative reasoning-action-observation loop. The workflow ends when the agent submits an answer or reaches the turn limit.

Benchmark Construction. We construct the dataset from thousands of open-source C/C++ projects. For each retained project, we build paired stripped and unstripped binaries: the stripped binaries form the benchmark inputs, while symbols and source mappings in the unstripped binaries identify compiled function entries and construct the hidden oracle. The validated pool contains 11,767 tasks across 529 projects.

We map source definitions to compiled entries and retain project-native functions with a clear purpose and meaningful, self-contained behavior, such as core algorithms, parsers, state transitions, data transformations, validation logic, resource management, and major request or command workflows. We exclude tests, bundled dependencies, thunks, aliases, compiler-generated variants, trivial utilities, placeholder functions, generic error handlers, and wrappers that delegate the described behavior. Task counts therefore vary by project. For each selected function, a generation model produces a behavior description from its source implementation and supporting evidence. The description may cover inputs, state changes, control-flow decisions, calls, constants, errors, and externally visible effects, but omits the function name, source location, address, and recognizable name variants. We reject incomplete records, leaked identifiers, and addresses that do not map to verified function entries, and deduplicate tasks using source-function, binary-address, and description similarity.

The reported evaluation samples 50 tasks from this larger pool. We randomly select five tasks from each of 10 representative projects—`tmux`, `tree`, `GNU Wget2`, `XZ Utils`, `YAJL`, `zlib`, `libuv`, `libgit2`, `nm`, and `Lua`. This project-balanced subset covers diverse software domains and function behaviors.

Scoring. A prediction is correct only when the submitted ELF virtual address exactly matches one of the accepted oracle entries. Every accepted alternative must be a verified function start that independently implements the complete described behavior; related callers, callees, file offsets, runtime addresses, and interior instructions are incorrect. A run receives zero credit if it submits an incorrect entry or fails to produce a valid submission, including non-convergent repeated tool use and runtime or API failures. We use the same Ghidra-backed interaction protocol and task budget for all compared models and compute accuracy over all 50 tasks, without excluding invalid submissions.

C.3 SWE-bench Verified Mini50 Harness

We evaluate software-engineering agent capability on a fixed internal 50-task subset of SWE-bench Verified (Jimenez et al., 2024; OpenAI, 2024). The evaluation has two stages. First, mini-SWE-agent v2 (SWE-agent, 2026; Yang et al., 2024) with tool calling performs an interactive repair trajectory inside each task’s Docker `/testbed` using a structured `bash` tool and submits a unified-diff patch. The official SWE-bench scorer then applies the patch and executes the task-specific `FAIL_TO_PASS` and `PASS_TO_PASS` tests. The evaluation uses a fixed Mini50 manifest, a temperature of zero, 10 shards of five tasks each, a 250-step interaction budget per task, and a 1,024-token output limit per turn. The ZGCM evaluation service uses a 256K-token vLLM (vLLM Team, 2023) context window with string-formatted tool observations, automatic tool choice, and GLM tool-call and reasoning parsers. Under this protocol, the best fully audited ZGCM-family evaluation resolves 2 of 50 tasks (4.0%).

C.4 Terminal-Bench 2.0 Harness

We evaluate terminal-agent capability on the standard 89-task Terminal-Bench 2.0 benchmark (Merrill et al., 2026) using task-specific Docker images and Harbor (Harbor Framework Team, 2026) for execution and verification. A thinking-enabled terminal agent using the GLM XML protocol interacts through a persistent `bash(command)` tool for at most 64 steps per task, with a 4,096-token output limit per turn, a temperature of zero, top- p of 0.95, a 180-second timeout per command, and a one-hour timeout per task. The vLLM inference service uses a 128K-token context window with the GLM-5.1 tool parser and GLM-4.5 reasoning parser. The 89 tasks are partitioned across four shards, with one attempt per task. Only tasks affected by predeclared infrastructure failures or verifier timeouts are rerun with identical parameters after cache warm-up. Under this protocol, the best fully audited ZGCM-family evaluation resolves 2 of 89 tasks (2.25%).

D Atomic Capability Evaluation Framework

Atomic Capability Evaluation (ACE) diagnoses model behavior at a finer granularity than aggregate benchmarks. The framework comprises 18 categories, 183 atomic capabilities, and 2,503 executable probes. This appendix specifies the taxonomy, scoring method, validation procedure, and category-level results.

D.1 Design Principles

ACE follows four design principles. **(P1) Atomicity.** Each capability has a narrow semantic boundary, so a failure identifies a specific behavior rather than only a broad domain. **(P2) Probe diversity.** Each capability is tested through multiple surface forms at three author-assigned difficulty levels (L1–L3), reducing sensitivity to any single wording. **(P3) Intent-aware scoring.** The evaluator applies strict response matching only when formatting is the target behavior; otherwise, it extracts and scores the substantive answer while normalizing non-semantic variation. **(P4) Traceability.** Every decision links to the probe, reference answer, scoring policy, extracted answer, and failure record.

D.2 Taxonomy and Scoring

Across 18 capability categories, ACE decomposes model behavior into 183 atomic capabilities. Each atomic capability targets a specific behavior and is evaluated with multiple probes that vary in surface form and difficulty. Table 13 summarizes the scope and allocation of this taxonomy. The 2,503 probes comprise 835 L1, 835 L2, and 833 L3 instances. These labels denote three levels of increasing probe difficulty.

Table 13 Evaluated ACE taxonomy, category scope, and probe allocation.

Category	Caps.	Probes	Primary evaluation scope
Agent	8	120	Environment interaction, long-horizon execution, recovery, memory, and coordinated action.
Arithmetic	13	135	Numerical operations, comparison, conversion, sorting, and exact calculation.
Causal	5	75	Mechanisms, interventions, counterfactuals, and latent-state inference.
Code	26	389	Program understanding, generation, debugging, optimization, testing, and artifact editing.
Dialogue	6	90	Multi-turn coherence, clarification, coreference, intent, persona, and subtext.
Instruction	13	195	Compliance with format, length, lexical, structural, stylistic, and multi-constraint requirements.
Knowledge	8	96	Recall and application of facts, formulas, APIs, conventions, and common mechanisms.
Language	27	321	Lexical, syntactic, semantic, multilingual, and character-level language processing.
Logic	10	102	Boolean, deductive, inductive, constraint, set, strategic, and structural reasoning.
Long Context	3	45	Position-robust retrieval, multi-span integration, and multi-document synthesis.
Mathematics	11	165	Algebra, discrete mathematics, limits, optimization, geometry, number theory, and proof.
Metacognition	6	90	Ambiguity and error detection, self-checking, uncertainty, and plan revision.
Output	7	105	Final-answer control, exact formatting, fluent prose, units, and structured-artifact generation.
Planning	6	90	Subgoal decomposition, forward and backward reasoning, state tracking, experiment design, and persistence.
Reading	11	141	Targeted retrieval, distractor filtering, passage integration, structured lookup, and cross-document reading.
Tool Use	7	105	Tool selection, argument extraction, schema compliance, invocation, and multi-tool coordination.
Truthfulness	8	119	Calibrated abstention, hallucination resistance, source fidelity, infeasibility detection, and misconception resistance.
World Modeling	8	120	Physical, spatial, temporal, social, procedural, and theory-of-mind reasoning.
Total	183	2,503	

Table 14 illustrates how these difficulty levels are instantiated for one atomic capability. The example is drawn from the *argument extraction* capability in the Tool Use category: L1 extracts a single scalar argument, L2 combines unit conversion with multiple arguments, and L3 constructs nested list and object arguments.

Table 14 Example L1–L3 probes for argument extraction.

Level	Example probe	Expected tool arguments
L1	<code>set_timer(minutes: int)</code> . “Start a timer for 45 minutes.”	<code>minutes=45</code>
L2	<code>ship_package(weight_grams: int, speed: standard/overnight)</code> . “Send a 2.5 kg parcel by overnight delivery.”	<code>weight_grams=2500,</code> <code>speed=overnight</code>
L3	<code>submit_order(items, shipping)</code> . “Order 3 of SKU A100 and 2 of SKU B250, and ship them to 12 Oak Lane in Canada.”	<code>items=[(A100,3),(B250,2)];</code> <code>shipping=(12 Oak Lane, Canada)</code>

ACE evaluates each probe with a scorer selected for the target capability. Among the 2,503 probes, 2,155 (86.1%) use deterministic scoring. Numeric scorers recognize answer markers, final lines, boxed expressions, units, and equivalent fractions. String scorers accept quoted and sentence-embedded answers, while multiple-choice scorers restrict letter matching to audited answer positions. Strict full-response matching is reserved for tasks that explicitly test format, length, or the absence of extra text. Code probes use abstract-syntax-tree checks or isolated unit tests, and tool-use probes validate arguments against JSON Schema. For the remaining 348 probes (13.9%), whose correctness cannot be expressed through a reliable executable criterion, the automated pipeline invokes a separate model with a predefined binary rubric. Both scoring paths execute within the same automated pipeline; human review is limited to benchmark construction and offline validation. Probe outcomes are aggregated into atomic-capability scores with L1–L3 breakdowns; these scores localize improvements and regressions, while category-level aggregation summarizes broader patterns.

D.3 AI-Assisted Benchmark Construction

We build the benchmark through iterative human-LLM co-design. At the taxonomy level, LLMs propose decompositions, boundary cases, and missing dimensions; researchers consolidate these proposals into capabilities with explicit evaluation objectives. At the probe level, LLMs draft questions, reference answers, difficulty variants, and adversarial cases. Researchers verify that each item isolates its intended capability and revise ambiguous or confounded items. At the scoring level, LLMs serve two roles. First, they translate evaluation specifications into candidate answer extractors, validators, static analyses, and task-specific tests; researchers validate and freeze these components before evaluation. Second, at run time, an LLM serves as the automated judge for rubric-based probes.

D.4 Evaluation Results

Figure 17 compares ZGCM-7B with the reference models across all 18 capability categories.

	Agent	Arith	Causal	Code	Dialog	Instr	Know	Lang	Logic	LongCtx	Math	Meta	Output	Plan	Read	Tool	Truth	World
Kimi K3	100.0	97.8	100.0	95.4	100.0	98.5	99.0	98.4	100.0	100.0	98.8	94.4	98.1	100.0	100.0	99.0	98.3	99.2
DeepSeek V4 Flash	99.2	97.0	100.0	93.8	100.0	90.8	100.0	98.4	99.0	100.0	100.0	95.6	99.0	97.8	100.0	100.0	91.6	99.2
Claude Opus 4.6	98.3	96.3	100.0	95.6	97.8	94.4	100.0	97.2	100.0	73.3	99.4	93.3	100.0	98.9	98.6	99.0	96.6	97.5
Claude Opus 4.7	100.0	92.6	100.0	96.4	98.9	91.8	100.0	95.3	99.0	80.0	100.0	94.4	90.5	98.9	98.6	99.0	100.0	99.2
Claude Sonnet 5	98.3	95.6	100.0	94.3	98.9	91.3	100.0	96.9	96.1	84.4	98.2	91.1	100.0	100.0	99.3	98.1	96.6	100.0
Claude Opus 4.8	99.2	94.8	100.0	95.1	98.9	89.7	100.0	96.3	99.0	88.9	98.8	92.2	93.3	97.8	99.3	98.1	97.5	99.2
DeepSeek V4 Pro	100.0	94.1	100.0	93.1	100.0	83.1	97.9	94.4	100.0	95.6	99.4	95.6	95.2	97.8	97.9	99.0	92.4	100.0
GPT-4o	96.7	80.0	100.0	88.2	100.0	73.8	97.9	85.0	91.2	88.9	87.9	88.9	92.4	83.3	92.9	96.2	91.6	92.5
ZGCM-7B (ours)	92.5	91.8	96.0	72.2	92.2	49.7	84.4	79.8	96.1	88.9	98.2	83.3	85.7	97.8	95.0	88.6	68.9	90.8
Qwen3-8B-256K	88.3	71.1	97.3	72.2	93.3	50.8	90.6	74.1	92.2	57.8	95.8	76.7	86.7	83.3	90.8	90.5	69.8	80.8

Figure 17 Category-level ACE performance across models. Rows denote models and columns denote the 18 evaluated categories; each cell reports the corresponding pass rate (%).

ZGCM-7B performs best in mathematics (98.18%), planning (97.78%), logic (96.08%), causal reasoning (96.00%), and reading (95.04%). Its weakest categories are instruction following (49.74%), truthfulness (68.91%), code (72.24%), and language (79.75%). The heatmap also shows that its gap to the strongest reference models is concentrated in a small set of categories rather than distributed across the capability space.

Compared with Qwen3-8B-256K, a model of similar scale, ZGCM-7B shows its largest advantages in long-context processing (+31.11 percentage points), arithmetic (+20.74), planning (+14.45), and world modeling (+10.00).

D.5 Quality and Scope

The benchmark was refined through repeated scoring audits, cross-model replay, regression testing, and an independent cross-check. These checks exposed ambiguities and scoring errors that were repaired before subsequent evaluation. This refinement process was not fully autonomous: human experts were required to adjudicate disagreements, verify capability boundaries, inspect systematic failure patterns, and approve substantive changes to probes and scorers. In practice, purely model-driven iteration was insufficient because it could reproduce earlier assumptions or overlook correlated errors; effective benchmark evolution therefore depended on continued human guidance and oversight.

ACE measures a structured set of atomic capabilities rather than complete end-to-end application performance. Its results depend on probe coverage and scoring quality, and rubric-based LLM judging may retain residual bias. We therefore use ACE as a diagnostic complement to domain benchmarks and application-level evaluation, rather than as a single comprehensive measure of model quality.

E AI Autonomy Rubric and Contributor Assessments

This appendix documents the task-specific autonomy rubric and individual contributor ratings underlying Section 6.2. The assessment covers 11 task categories grouped into seven R&D domains. Data engineering comprises data cleaning, data acquisition, and synthetic data generation. Infrastructure engineering comprises operator design, training and inference framework design, and resource platform management. The remaining

domains are model architecture design, learning algorithm design, experimentation and monitoring, model evaluation, and deployment engineering.

E.1 Assessment Scope and Interpretation

Nine core contributors each provided one rating for every task category, yielding 99 ratings. Contributors are represented by numeric identifiers in the reporting table, with each identifier referring to the same contributor across all tasks. The assessments capture contributor judgments and should not be interpreted as standardized capability measurements.

The rubric distinguishes five levels according to the allocation of responsibility between humans and AI, the ability to adapt workflows, and the need for human intervention. L1 describes assistance with individual steps; L2 describes execution of predefined workflows; L3 introduces adaptive execution with human approval at consequential decision points; and L4 describes autonomous iteration within human-defined objectives and constraints. L5 extends autonomy to identifying research objectives and coordinating sustained work across R&D stages.

These levels characterize autonomy rather than productivity, output quality, or scientific novelty. Higher autonomy does not necessarily imply better results, and substantial benefits may arise from assistance at lower levels. L2 behavior may also be implemented through conventional automation; the relevant distinction in an AI assessment is the role played by the agent within that workflow.

L5 serves as a reference for full autonomy, rather than an assertion of demonstrated capability. It does not imply independence from prior knowledge, freedom from operating constraints, or the elimination of external validation. Its defining distinction from L4 is autonomous objective formation and sustained coordination across stages.

The task-specific descriptions below characterize representative behavior at each level. Autonomy demonstrated in a narrow subtask should not be interpreted as equivalent autonomy across an entire task category.

E.2 Task-Specific Autonomy Criteria

E.2.1 Data Cleaning

- L1: Basic Assistance.** Humans define cleaning rules and inspect outputs. AI assists with individual regular expressions, formatting operations, or local code corrections.
- L2: Partial Automation.** AI executes predefined cleaning pipelines, including filtering, deduplication, sensitive-data removal, and fixed-format quality reporting.
- L3: Conditional Autonomy.** AI identifies distributional anomalies and low-quality samples, proposes or adapts cleaning strategies, and requests approval before consequential changes to data selection or processing.
- L4: High Autonomy.** Given downstream objectives and constraints, AI develops operational quality criteria and iteratively cleans, selects, and mixes data, validating decisions against downstream outcomes.
- L5: Full Autonomy.** AI identifies evolving data-quality needs and continuously revises cleaning strategies in coordination with acquisition, synthesis, and training, without routine human direction.

E.2.2 Data Acquisition

- L1: Basic Assistance.** Humans identify, select, and integrate data sources. AI assists with source discovery, download scripts, or interface documentation.
- L2: Partial Automation.** AI acquires, parses, and stores data using predefined source lists, interfaces, and collection rules.
- L3: Conditional Autonomy.** AI identifies coverage gaps, discovers candidate sources, and adjusts collection plans. Sources requiring additional authorization or consequential spending are referred to humans.

- L4: High Autonomy.** Given coverage objectives and access constraints, AI independently discovers, evaluates, integrates, and refreshes sources, adjusting acquisition strategies based on observed gaps.
- L5: Full Autonomy.** AI identifies emerging data needs and continuously evolves acquisition workflows in coordination with other R&D stages, while operating within established access and usage constraints.

E.2.3 Synthetic Data Generation

- L1: Basic Assistance.** Humans design generation templates and quality criteria. AI produces samples in response to individual instructions, with humans selecting the resulting data.
- L2: Partial Automation.** AI generates data at scale using fixed prompts, templates, or generators, followed by predefined filtering and deduplication.
- L3: Conditional Autonomy.** AI selects generation methods based on coverage gaps and iterates on prompts, filters, and data mixtures. Humans approve consequential changes to quality criteria or generation objectives.
- L4: High Autonomy.** Given training objectives, AI autonomously generates, evaluates, filters, and resamples synthetic data, using downstream results to guide successive iterations.
- L5: Full Autonomy.** AI identifies new synthesis objectives and develops generation strategies in coordination with model, algorithm, and evaluation changes, sustaining the process without routine human direction.

E.2.4 Model Architecture Design

- L1: Basic Assistance.** Humans determine model structure, scale, and module configuration. AI provides reference information, code completion, or local implementation suggestions.
- L2: Partial Automation.** AI instantiates established architecture templates, generates configurations and implementation code, and runs predefined architectural variants.
- L3: Conditional Autonomy.** AI proposes, implements, and compares architecture variants based on objectives and experimental results. Major structural or scale changes require human approval.
- L4: High Autonomy.** Given capability targets and resource constraints, AI autonomously searches, implements, trains, and validates architectures through multiple iterations.
- L5: Full Autonomy.** AI identifies architectural research directions and develops model structures in coordination with data, learning algorithms, and infrastructure, without routine human direction.

E.2.5 Learning Algorithm Design

- L1: Basic Assistance.** Humans design training objectives, losses, optimizers, and reinforcement learning methods. AI assists with implementation and local debugging.
- L2: Partial Automation.** AI configures established algorithm templates and evaluates predefined combinations of losses, optimizers, sampling strategies, or reinforcement learning procedures.
- L3: Conditional Autonomy.** AI analyzes training results and proposes algorithmic changes. Consequential changes to learning objectives or training methods require human approval.
- L4: High Autonomy.** Given learning objectives and constraints, AI independently designs, implements, evaluates, and refines learning algorithms across multiple experimental iterations.
- L5: Full Autonomy.** AI identifies algorithmic research questions, develops candidate methods and supporting analyses, and coordinates their validation with other R&D stages without routine human direction.

E.2.6 Experimentation and Monitoring

- L1: Basic Assistance.** Humans formulate hypotheses, design and launch experiments, and monitor progress. AI assists with code completion or interpretation of individual results, status messages, and errors.
- L2: Partial Automation.** AI launches experiments from fixed templates, runs predefined parameter sweeps, and applies rule-based monitoring, alerts, and retries.
- L3: Conditional Autonomy.** AI designs routine experiments, monitors execution, diagnoses anomalies, and adapts the experimental plan. Major failures, budget changes, and consequential interpretations are referred to humans.
- L4: High Autonomy.** Given research objectives and resource limits, AI autonomously plans, executes, monitors, troubleshoots, and analyzes successive experiments, maintaining records of evidence and conclusions.
- L5: Full Autonomy.** AI identifies research questions, formulates hypotheses, allocates resources within operating constraints, and uses experimental conclusions to determine subsequent research directions.

E.2.7 Operator Design

- L1: Basic Assistance.** Humans design, implement, and debug operators. AI provides code completion, interface explanations, or local performance suggestions.
- L2: Partial Automation.** AI generates or modifies operators using established templates and runs fixed correctness and performance benchmarks.
- L3: Conditional Autonomy.** AI identifies bottlenecks, proposes specialized operators, and implements and validates candidate solutions. Consequential hardware-specific changes require human approval.
- L4: High Autonomy.** Given correctness and performance targets, AI autonomously designs, implements, compiles, tunes, and repairs operators across the specified hardware environments.
- L5: Full Autonomy.** AI identifies operator and hardware co-design opportunities and evolves implementation and optimization strategies alongside changes in models, compilers, and hardware.

E.2.8 Training and Inference Framework Design

- L1: Basic Assistance.** Humans construct frameworks and connect their components. AI assists with configuration, integration code, and error interpretation.
- L2: Partial Automation.** AI configures standard frameworks and pipelines, executing predefined builds, integration tests, and task orchestration.
- L3: Conditional Autonomy.** AI adapts parallelization strategies and runtime configurations to workloads and diagnoses routine framework problems. Major framework changes require human approval.
- L4: High Autonomy.** Given training and inference objectives, AI independently improves framework components and validates changes to parallel execution, compilation, runtime behavior, and pipeline organization.
- L5: Full Autonomy.** AI identifies system design objectives and continuously restructures and validates frameworks across model and hardware changes, coordinating improvements with the broader R&D workflow.

E.2.9 Resource Platform Management

- L1: Basic Assistance.** Humans request and allocate resources, submit jobs, and maintain environments. AI suggests commands or explains individual failures.
- L2: Partial Automation.** AI executes predefined resource workflows, including environment setup, queuing, job retries, and rule-based preemption.

- L3: Conditional Autonomy.** AI dynamically balances workloads, isolates unhealthy nodes, and migrates jobs. Actions with substantial cross-tenant impact or resource implications require human approval.
- L4: High Autonomy.** Given workload objectives and resource limits, AI forecasts demand, adjusts capacity, recovers from failures, and optimizes placement, communication, and cluster utilization.
- L5: Full Autonomy.** AI identifies platform-level improvement objectives and coordinates scheduling and configuration across clusters and hardware types, continuously adapting the platform to evolving R&D needs.

E.2.10 Model Evaluation

- L1: Basic Assistance.** Humans select benchmarks, design evaluation procedures, and inspect results. AI assists with individual scripts, metric explanations, or local analysis.
- L2: Partial Automation.** AI runs predefined evaluation pipelines and produces metric comparisons, visualizations, and structured reports.
- L3: Conditional Autonomy.** AI adaptively investigates weaknesses, clusters failure cases, and conducts follow-up diagnostics. Changes to evaluation objectives or consequential judgments require human review.
- L4: High Autonomy.** Given evaluation objectives, AI designs and validates additional tests, conducts adversarial evaluations, and iteratively investigates model weaknesses, including checks for benchmark-specific overfitting.
- L5: Full Autonomy.** AI identifies emerging evaluation needs and continuously develops evaluation criteria and benchmarks in coordination with model development, maintaining validation appropriate to each claim.

E.2.11 Deployment Engineering

- L1: Basic Assistance.** Humans package models, configure environments, deploy services, scale capacity, and perform rollbacks. AI provides scripts or suggestions for individual operations.
- L2: Partial Automation.** AI executes predefined continuous integration and delivery pipelines, serving configurations, health checks, and rollback procedures.
- L3: Conditional Autonomy.** AI selects quantization settings, instance configurations, and rollout strategies based on deployment objectives, and manages staged releases and monitoring. Full releases and consequential rollbacks require human approval.
- L4: High Autonomy.** Given quality, latency, throughput, and cost targets, AI independently performs model compression, quantization, serving optimization, staged deployment, scaling, and rollback within authorized limits.
- L5: Full Autonomy.** AI identifies evolving deployment objectives and continuously adapts the deployment and inference system to model, traffic, and environment changes, coordinating with training, evaluation, and resource management.

E.3 Individual Contributor Ratings

Table 15 reports the complete contributor ratings. Tasks are presented in the same order as in Figure 14, and columns identify Core Contributors 1–9. The table preserves individual variation that may be obscured by aggregate summaries.

For each task, the figure reports the arithmetic mean and sample standard deviation of the nine ratings after encoding L1–L5 as 1–5. These statistics provide descriptive summaries of ordinal responses; they do not establish equal capability differences between successive levels. Error bars represent variation across contributors rather than confidence intervals. The individual points in the figure correspond directly to the entries below.

Table 15 Individual assessments of AI autonomy across R&D tasks. Columns 1–9 denote Core Contributors 1–9, with identifiers held constant across tasks. Ratings follow the task-specific criteria in Appendix E.2.

Task	Core Contributor								
	1	2	3	4	5	6	7	8	9
Data cleaning	L3	L2	L3	L3	L2	L3	L3	L2	L2
Data acquisition	L4	L4	L4	L4	L1	L2	L4	L2	L4
Synthetic data generation	L4	L3	L3	L3	L3	L3	L3	L3	L3
Model architecture design	L3	L2	L2	L2	L1	L2	L2	L1	L2
Learning algorithm design	L3	L2	L2	L3	L2	L1	L2	L1	L2
Experimentation & monitoring	L5	L4	L4	L4	L4	L4	L4	L4	L3
Operator design	L4	L4	L4	L3	L4	L3	L3	L4	L4
Training & inference framework design	L3	L3	L3	L2	L2	L3	L2	L3	L2
Resource platform management	L4	L4	L3	L4	L2	L4	L4	L4	L3
Model evaluation	L4	L4	L4	L4	L1	L4	L4	L4	L3
Deployment engineering	L4	L4	L4	L4	L2	L4	L4	L4	L4

F Evaluation Reproducibility and Provenance

Reproducibility requires binding a numerical result to four layers at once: the model artifact, the evaluation data, the invocation contract, and the scoring implementation. For each final benchmark, the evaluation record should preserve the model hash, official dataset URL and immutable revision, split, prompt wrapper, few-shot examples, system prompt, chat template, mode-control setting, answer extraction, metric implementation, random seed, timeout, inference engine and revision, hardware, and raw-output location. External comparison values are taken from the cited releases rather than a unified rerun.

We evaluate ZGCM-1 in no-think mode using deterministic decoding with temperature 0, top- $p = 1.0$, and one run.

Table 16 Selected no-think benchmark results (%) for ZGCM-1-7B and selected 7B–8B models. Dark blue cells in bold indicate the best result; light blue cells indicate the second-best result.

Benchmark	ZGCM-1-7B	Olmo 3 7B Instruct	Qwen3-8B (non-thinking mode)	Qwen2.5-7B
<i>Reasoning & General</i>				
MATH-500	89.40	87.30	82.30	71.00
AIME 2024	33.33	44.30	26.20	11.30
AIME 2025	43.33	32.50	21.70	6.30
BBH	69.48	71.20	73.70	68.80
<i>Code</i>				
HumanEval+	66.46	77.20	79.80	74.90
MBPP+	75.93	60.20	64.40	62.60
LiveCodeBench v3	34.31	29.50	53.20	34.50
<i>Knowledge</i>				
MMLU	70.03	69.10	80.40	77.20
GPQA-Diamond	42.42	40.40	44.60	35.60
<i>Instruction Following</i>				
IFEval	78.00	85.60	86.30	73.40
IFBench (Pyatkin et al., 2025)	31.67	32.30	29.30	28.40